A Predictive Maintenance Platform for Agriculture Tractors in Rwanda

BSc. in Software Engineering

Josiane Ishimwe

Machine Learning & Artificial Intelligence

Supervisor:

Marvin Ogore Muyonga

Date: 17 November 2025

DECLARATION

This Proposal Project is my original work, unless stated, and all external sources have been referenced or cited in my document. This work has not been presented for the award of a degree or for any similar purpose in any other university.

Signature Date: 17 November 2025

Name of Student Josiane Ishimwe

CERTIFICATION

The undersigned certifies that he has read and hereby recommends for acceptance by African Leadership University a report entitled ………..

Signature……………………… Date………………………….

Prof/Dr./Mrs./Miss/Mr. Name of the Supervisor

Faculty,

Bachelor of Software Engineering,

ALU

DEDICATION AND ACKNOWLEDGEMENT

This work is dedicated to my family and to all smallholder farmers in Rwanda, whose resilience and commitment continue to feed our communities. Your contribution is irreplaceable. It is time to stop romanticising agriculture and begin professionalising it, because agriculture is not guesswork, and neither is mechanisation. May this work, in a small way, support that transformation.

First and foremost, I give all glory and thanks to God for the gift of life, strength, and the grace to complete this undergraduate journey. Every step, seen and unseen, was possible because of You.

My sincere appreciation goes to my Capstone supervisor and Machine Learning coach, Mr. Marvin Ogore, for his guidance, patience, and support. My gratitude also extends to all BSE facilitators who shaped my learning.

To my family, this is not for me, but for us. Thank you for your love, support, and care. And to my mother especially: thank you for every prayer, every word of encouragement, and every silent sacrifice. Your support has carried me farther than words can ever express.

To everyone from Hello Tractor, Software Engineering Manager Bello Moussa Amadou, Customer Success Manager David Kayi, Product Designer Delight Asaph, and Kayonza Hub Manager Mutabazi Yassin and to the RICA technician and students who supported me throughout this research: every piece of advice, every suggestion, every feedback, and every moment of help meant more than you know. Thank you.

To my classmates and cohortmates who became more than friends, thank you for the shared struggles, the laughter, and the sense of belonging that made this journey lighter and unforgettable.

Last but not least, I want to thank myself. Thank you for not giving up, for showing up even on the days when my soul was tired. For choosing to push forward, to believe, and to trust that the finish line was worth it. You made it.

Abstract

As Rwanda advances toward transforming its agricultural sector into a profitable and high-value professional industry, mechanisation has emerged as a core strategy. However, tractor adoption remains low at only 0.8% due to high equipment costs, poor maintenance practices that shorten the tractor's lifespan, and frequent unexpected breakdowns, causing 20–30% downtime and $50–100 in avoidable annual repairs. Although government programs and private initiatives are working to reduce the financial barrier to tractor ownership, smallholders still struggle to access affordable maintenance technologies. Existing predictive maintenance systems are often unsuitable because they require costly annual subscriptions of $500, depend on stable internet connectivity, and rely on hardware incompatible with older tractor models. This research developed TractorCare, a smartphone-based platform that combines machine learning acoustic analysis to enable early problem detection without expensive hardware, along with rule-based maintenance scheduling to digitise maintenance service scheduling using manufacturer manuals. A ResNet-like CNN trained on tractor sound recordings achieved 81% accuracy with 100% recall, ensuring zero missed mechanical problems. Field testing in Kayonza District validated diagnostic capability through 99% confidence detection of genuinely problematic tractors, while personalised baseline monitoring distinguished individual tractor acoustic signatures from mechanical deterioration. The system digitises Massey Ferguson 240 and 375 maintenance manuals, automatically scheduling service tasks based on engine hours tracked through daily usage logs. Pilot deployment with 10 farmers demonstrated 96% offline functionality, 18-25 second response times, strong acceptance, 7.7/10 satisfaction, and 100% retention. Results validate that acoustic analysis enables transition from reactive to proactive maintenance, offering Rwanda's smallholders previously inaccessible predictive capabilities supporting PSTA5 and Vision 2050 agricultural mechanisation goals.

Table of Contents

DECLARATION 2

CERTIFICATION 3

DEDICATION AND ACKNOWLEDGEMENT 4

Abstract 5

Table of Contents 6

List of Tables 9

List of Figures 10

List of Acronyms/Abbreviations 11

CHAPTER ONE: INTRODUCTION 12

1.1 Introduction and Background 12

1.2 Problem statement 13

1.3 Project's main objective 14

1.3.1 List of the specific objectives 15

1.4 Research questions 16

1.5 Project scope 16

1.6 Significance and Justification 16

1.7 Research Budget 17

1.8 Research Timeline 18

1.9: Ethical Considerations 18

CHAPTER TWO: LITERATURE REVIEW 21

2.1 Introduction 21

2.2 Historical Background of the Research Topic 21

2.3 Overview of Existing System And Related Work 22

2.3.1 IoT and Sensor Integration in PdM Software (The Commercial Standard) 22

2.3.2 Machine Learning Algorithms for Failure Prediction 23

2.3.2.2 Traditional Machine Learning Models 23

2.3.2.2 Deep Learning Models 23

2.3.3 The Research Gap: Audio-Based Condition Monitoring 24

2.4 Summary of Reviewed Literature.. 24
2.5 Strengths and Weaknesses of the Existing System(s).. 30
2.6 General Comments...31
CHAPTER THREE: SYSTEM ANALYSIS AND DESIGN..33
3.1 Introduction.. 33
3.2 Research Design (including the SDLC model used)..34
3.2.1 Dataset and Dataset Description.. 36
3.3 Functional and Non-functional Requirements.. 36
3.3.1 Functional Requirements (FR)..36
3.3.2 Non-Functional Requirements (NFR)...39
3.2.1 Machine Learning Model Diagram.. 42
3.5 Float Chart, Use Case Diagram, Sequence Diagram and all other diagrams...............44
3.5.1 Float Chart...45
3.5.2 Use Case Diagram:...46
3.5.3 Sequence Diagram..47
3.5.4 ERD Diagram...48
CHAPTER 4: SYSTEM IMPLEMENTATION AND TESTING....................................50
4.1 Implementation and Coding.. 50
4.1.1 Introduction... 50
4.1.2 Description of Implementation: Tools and Technology.......................................50
4.2 Graphical view of the project...52
4.2.1 Screenshots with description...52
4.2.1.1 Authentication and Onboarding.. 53
4.2.2 Key Code Implementations...63
4.3 Testing... 75
4.3.1 Introduction.. 75
4.3.2 Objective of Testing... 75
4.3.3 Unit Testing Outputs.. 75
4.3.4 Functional Validation Results...76

4.3.5 Integration Testing Results.. 78
4.3.6 System Performance Testing.. 79
4.3.7 Field Testing and User Acceptance.. 79
4.3.7.1 Testing Protocol.. 80
4.3.7.2 Testing Approach.. 80
4.3.7.3 Key Findings.. 81
CHAPTER 5: RESULTS AND SYSTEM DESCRIPTION.. 84
5.1 Introduction.. 84
5.2 System Deployment.. 84
5.2.1 Deployment Overview.. 84
5.2.2 Deployment Steps.. 84
5.2.3 Field Installation.. 84
5.3 User Guide and System Operation.. 85
5.3.1 Getting Started.. 85
5.3.2 Recording Engine Sounds.. 85
5.3.3 Predictions and Alerts.. 85
5.3.4 Maintenance Scheduling.. 85
5.3.5 Troubleshooting.. 85
5.4 Source Code and System Access.. 85
5.5 Machine Learning Model Performance.. 86
5.5.1 Foundation Training on Industrial Dataset.. 86
Table 5.1: Foundation Model Performance Comparison on MIMII Dataset.................... 87
5.2.2 Transfer Learning to Real Tractor Sounds.. 88
5.3 Field Testing and Real-World Validation.. 90
5.3.1 Field Testing Protocol and Environment.. 91
5.3.4 Field Testing Limitations and Future Validation Needs.. 91
5.4 System Performance and Responsiveness.. 92
5.4.1 Response Time Analysis.. 92
5.5 Offline Capability and Data Synchronisation.. 92

5.5.1 Offline Feature Coverage.......... 92
5.5.2 Synchronisation Performance and Reliability.......... 93
5.6 User Acceptance and Stakeholder Feedback.......... 93
5.6.1 Pre-Deployment Stakeholder Survey.......... 93
5.6.2 Pilot Program Deployment Results.......... 94
5.6.3 Usability Feedback and User Experience.......... 95
5.7 Discussion and Interpretation.......... 95
5.7.1 Achievement of Research Objectives.......... 95
5.7.2 Comparison with Existing Maintenance Approaches.......... 96
5.6.4 Comparative Performance Table.......... 97
5.7.3 Implications for Rwandan Agricultural Mechanisation.......... 98
CHAPTER 6: CONCLUSIONS AND RECOMMENDATIONS.......... 99
6.1 Discussion and Conclusions.......... 99
6.2 Advantages of the System.......... 100
6.3 Limitations of the Study.......... 101
6.4 Recommendations.......... 102
6.5 Suggestions for Further Studies or Research.......... 103
6.6 Final Conclusion.......... 104
References.......... 105

List of Tables

Table 1.1: Research Budget

Table 3.1: Functional Requirements

Table 3.2: Non-Functional Requirements

Table 3.3: Development Tools

Table 4.1: Backend Unit Testing Results

Table 4.2: Functional Requirements Validation Results

Table 4.3: Integration Testing Results

Table 4.4: System Performance Testing Results

Table 4.5: Field Testing Results Summary

Table 5.1: Foundation Model Performance Comparison on MIMII Dataset

Table 5.2: Comparative Performance Table

List of Figures

Figure 1.1: Research Timeline (Gantt-style Breakdown for 3 Months)

Figure 3.1: ML Model Architecture Diagram

Figure 3.2: System Architecture Diagram

Figure 3.3: Flow Chart

Figure 3.4: Use Case Diagram

Figure 3.5: Sequence Diagram

Figure 3.6: ERD Diagram

Figure 4.1: User Registration Screen

Figure 4.2: Login Screen

Figure 4.3: Home Dashboard3

Figure 4.4: Tractor Fleet Management Screen

Figure 4.5: Add Tractor Registration Form

Figure 4.6: Tractor Detail Screen - Good Health Status

Figure 4.7: Baseline Collection Workflow Progression

Figure 4.8: Audio Recording Interface with Recent Predictions

Figure 4.9: Maintenance Scheduling Interface

Figure 4.10: Usage History Logging Interface

Figure 4.11: Baseline Validation - Same Tractor vs. Different Tractor

Figure 4.12: Critical Anomaly Detection

Figure 5.1: [Model Comparison Chart - if included]

Figure 5.2:Models Comparison

Figure 5.3: Transfer Learning Performance Metrics

Figure 5.4: Transfer Learning Confusion Matrix

Figure 5.5: Field Testing Classification Accuracy by Tractor Condition

Figure 5.6: Stakeholder Interest in Predictive Maintenance Technology

Figure 5.7: Current Tractor Monitoring Practices

List of Acronyms/Abbreviations

AES - Advanced Encryption Standard

AI - Artificial Intelligence

API - Application Programming Interface

APK - Android Package Kit

CNN - Convolutional Neural Network

CORS - Cross-Origin Resource Sharing

CPU - Central Processing Unit

CRUD - Create, Read, Update, Delete

CSV - Comma-Separated Values

Dart - Programming language used with Flutter

FAO - Food and Agriculture Organisation

FastAPI - Modern Python web framework

FFT - Fast Fourier Transform

Flutter - Cross-platform mobile development framework

FR - Functional Requirement

FY - Fiscal Year

GDP - Gross Domestic Product

GPRS - General Packet Radio Service

GSM - Global System for Mobile Communications

HTML - HyperText Markup Language

HTTP - HyperText Transfer Protocol

HTTPS - HyperText Transfer Protocol Secure

IFPRI - International Food Policy Research Institute

IoT - Internet of Things

iOS - iPhone Operating System

JSON - JavaScript Object Notation

JWT - JSON Web Token

KNN - K-Nearest Neighbours

LSTM - Long Short-Term Memory

MFCC - Mel-Frequency Cepstral Coefficient

MIMII - Malfunctioning Industrial Machine Investigation and Inspection

MINAGRI - Ministry of Agriculture and Animal Resources (Rwanda)

ML - Machine Learning

MongoDB - NoSQL database system

NAEB - National Agricultural Export Development Board (Rwanda)

NFR - Non-Functional Requirement

NISR - National Institute of Statistics of Rwanda

NST - National Strategy for Transformation

ODM - Object-Document Mapper

PAYG - Pay-As-You-Go

PdM - Predictive Maintenance

PDX - Predictive Diagnostics (Predictronics software)

PSTA - Strategic Plan for Agriculture Transformation

Python - High-level programming language

React - JavaScript library for building user interfaces

ResNet - Residual Neural Network

RF - Random Forest

ROC-AUC - Receiver Operating Characteristic - Area Under Curve

RWF - Rwandan Franc

SDLC - Software Development Life Cycle

SEO - Search Engine Optimisation

SHAP - Shapley Additive Explanations

SNS - Smart Nkunganire System

SQLite - Lightweight relational database

SSA - Sub-Saharan Africa

SVM - Support Vector Machine

TensorFlow - Open-source machine learning framework

TLS - Transport Layer Security

UI - User Interface

USD - United States Dollar

UX - User Experience

VGG - Visual Geometry Group (CNN architecture)

VM - Virtual Machine

WAV - Waveform Audio File Format

CHAPTER ONE: INTRODUCTION

1.1 Introduction and Background

Agriculture serves as the backbone of Rwanda's economy, employing approximately 64.5% of the population and contributing nearly 25% to the national Gross Domestic Product (GDP) (NISR, 2023). Smallholder farmers, who cultivate about 70% of the country's agricultural land on fragmented plots averaging less than one hectare, play a critical role in ensuring food security and sustaining rural livelihoods (Philipp, 2022). Despite the challenges posed by Rwanda's rugged terrain and variable climate, the sector remains resilient, with crop production accounting for approximately 70% of agricultural output (FAO, 2022). Agricultural exports reached approximately USD 1.0 billion in FY 2023/24, projected to USD 1.2 billion in 2025, reflecting the sector's growing economic significance amid efforts to shift from subsistence to commercial farming (NAEB, 2025; World Bank, 2025). Between 2017 and 2023, increased yields and market access enabled agriculture to lift over 1 million people out of poverty, underscoring its pivotal role in poverty reduction (World Bank, 2024).

This potential, coupled with agriculture's role as the primary source of food security and a key driver of rural livelihoods, has prompted the Rwandan government to prioritise agricultural modernisation through targeted policies and programs. For instance, the Nkunganire program, a national agro-input subsidy initiative, aims to boost farm output and farmers' incomes by providing subsidised fertilisers, improved seeds, and other inputs. It is managed via the Smart Nkunganire System (SNS), a digital platform that digitises the entire supply chain, offers farmers mobile access to advisory services and financial tools, and ensures transparency and efficiency in input distribution and payments (MINAGRI, 2024a). The National Strategy for Transformation (NST1, 2017–2024) and its successor, NST2 (2024–2029), target a 6% annual increase in agricultural productivity through subsidies for fertilisers and improved seeds (MINAGRI, 2024b). Similarly, the Strategic Plan for Agriculture Transformation (PSTA4, 2018–2024) and the forthcoming PSTA5 emphasise land consolidation, expanding irrigation to 100,000 hectares by 2027, and developing export-oriented value chains (MINAGRI, 2025). Vision 2050 envisions transforming Rwanda into an upper-middle-income country by 2035 and a high-income country by 2050, with agriculture evolving into a high-value, export-driven sector contributing 30% of GDP through agro-processing and climate-smart practices. This vision supports the transition of smallholder farmers to commercial farming, aiming for an annual income of above USD 1,000 per capita (Government of Rwanda, 2020).

Globally, agricultural mechanisation has proven to be a cornerstone of productivity enhancement. By replacing manual labour with tractors, harvesters, and other machinery, cultivation efficiency can increase by 50–100% and reduce drudgery for farmers, with effects most pronounced in developed regions of North America and Asia (FAO, 2022). Evidence from Asia, for example, suggests that tractor adoption is associated with a 40% increase in output per hectare (IFPRI, 2021). While these technologies are transforming agriculture worldwide, their adoption in Rwanda is shaped by structural and resource constraints. Compared to regional peers like Kenya, where approximately 2.5% of farmland is mechanised, or Ghana, where approximately 1.8% is mechanised, Rwanda lags at 0.8% mechanisation using mechanical equipment, including approximately 0.5-0.8 tractors per 100 km^2 arable land (NISR, 2025; IFPRI, 2024).

Current tractor management and maintenance practices in Rwanda rely heavily on traditional, reactive approaches. Breakdowns are typically addressed only after they occur, leading to prolonged downtime, high repair costs, and reduced productivity. Coordination of tractors often involves manual processes, such as in-person bookings or informal arrangements, which exacerbate inefficiencies in a sector already constrained by limited access to spare parts and skilled technicians (Khumalo, 2024).

1.2 Problem statement

In Rwanda, the adoption of agricultural mechanisation among smallholder farmers, who constitute approximately 70% of the farm workforce (Philipp, 2022), is gradually increasing, primarily through Pay-As-You-Go (PAYG) service providers and cooperatives. These models have played a pivotal role in making tractors and other mechanised equipment accessible to farmers who would otherwise be unable to afford direct ownership. By reducing the financial burden of purchasing and maintaining machinery, they have enabled smallholders to leverage modern equipment for improved productivity.

However, their impact remains limited by frequent equipment breakdowns, poor maintenance practices that reduce equipment lifespan, and the absence of a digital solution tailored to Rwanda's conditions. As a result, tractor owners experience 20–30% equipment downtime, annual repair costs of USD 50–100 per farmer, and a decline in confidence in mechanised farming (Khumalo, 2024). This challenge is particularly severe for farmers accessing tractors through PAYG schemes or loans, who face the dual burden of continuing loan payments while experiencing lost productivity from non-functional equipment, causing premature equipment failure well before expected operational lifespan. Consequently, only 0.8% of Rwanda's farmland is mechanised, well below Kenya (2.5%) and Ghana (1.8%) (IFPRI, 2024).

Globally, Machine Learning and Internet of Things (IoT) technologies have enhanced mechanisation through predictive maintenance and monitoring systems. Platforms such as MF Connect and Predictronics reduce downtime and improve reliability (Massey Ferguson, 2023; Predictronics, 2024). Yet, these tools remain inaccessible to Rwandan smallholders due to high subscription costs of approximately USD 500 per year, limited internet coverage in 40% of rural areas, and incompatibility with older tractor models most commonly used in Rwanda (Adoyi, 2025; Khumalo, 2024).

In a regional context, Hello Tractor made strides in improving tractor access across sub-Saharan Africa, including Rwanda, through a digital platform that facilitated fleet management and utilisation (Hello Tractor, 2024). While effective in enhancing access, Hello Tractor fell short in providing predictive maintenance or real-time equipment health monitoring. Recent expansions, such as mechanisation hubs in Kayonza (2024) and Nyagatare (2025), improved services but still emphasised reactive over predictive approaches. This gap left tractors vulnerable to unexpected failures, undermining the reliability of services, particularly in remote areas with limited access to spare parts and skilled technicians (Adoyi, 2025). Addressing these gaps requires a localised, affordable predictive maintenance solution that supports offline functionality, automates maintenance schedules, and works with older tractor models.

1.3 Project’s main objective

This project aimed to design and develop a hybrid predictive maintenance platform for Rwanda's smallholder tractor fleets by combining a rule-based maintenance system with audio-based anomaly detection. The platform works in two ways: Rule-Based Maintenance, which uses manufacturer manuals and guidelines to schedule regular maintenance; and Audio-Based Anomaly Detection, which uses a Convolutional Neural Network (CNN) to listen to tractor engine sounds and classify their condition as Good, Warning, or Critical.

A cross-platform mobile app was built using Flutter with offline functionality to support rural areas with poor internet. Farmers can manage multiple tractors, record engine sounds, and get automated maintenance alerts based on usage, engine hours, and manufacturer recommendations. The system tracks tasks and sends updates to ensure proactive maintenance is followed.

Overall, the project aimed to reduce unexpected breakdowns and repair costs, extend tractor lifespan, and improve efficiency by helping smallholder farmers follow preventive maintenance practices. This makes tractor use more reliable, predictable, and sustainable for small farms in Rwanda.

1.3.1 List of the specific objectives

1. To conduct comprehensive stakeholder research and requirements validation through consultations with 15–20 participants in Kayonza District. These included smallholder farmers, mechanics, PAYG providers, and cooperative managers. The study documented existing maintenance practices, cost structures, and operational challenges using structured surveys and interviews. Insights from this phase informed the user-centred design of the mobile application, ensuring that its functionality aligned with real user needs and local operating conditions, particularly regarding maintenance literacy. This activity was completed by September 2025.

2. To design and develop a robust offline-to-online synchronisation mobile application with the following capabilities: (1) ML-based audio classification achieving greater than 80% accuracy for normal/abnormal engine sound detection, (2) rule-based maintenance scheduling achieving greater than 90% accuracy based on manufacturer guidelines and usage intensity, and (3) an informational website. Development and initial testing were completed by October 2025.

3. To evaluate the effectiveness of the developed platform through technical, user, and economic validation. Technical validation assessed machine learning model performance using metrics such as precision, recall, and F1-score, alongside an evaluation of rule-based schedule accuracy. User validation involved usability testing with selected stakeholders to determine ease of use and improvements in maintenance awareness and adherence to scheduling behaviour. Finally, an economic analysis assessed the platform's potential to reduce downtime by at least 20% and enhance maintenance tracking and adoption among

farmers, cooperatives, and PAYG providers. This evaluation phase was completed by November 2025.

1.4 Research questions

1. What are the specific, localised maintenance and accessibility barriers (cost, technical literacy, repair time, spare part availability) faced by tractor owners in the Rwandan agricultural sector, and how do these factors impact tractor downtime and maintenance adherence among operators?
2. How do current commercial and regional Predictive Maintenance solutions fail to meet the economic and infrastructural needs of Rwanda's smallholder mechanisation context?
3. What design and technical factors are essential for creating an affordable, predictive maintenance and management platform suited to Rwanda's smallholder mechanisation context?
4. How effective are machine learning models, trained on localised tractor data, in anticipating maintenance needs and minimising disruptions for commonly used tractors in Rwanda?
5. To what extent can the proposed platform enhance tractor reliability through reduction in unscheduled downtime and improve maintenance efficiency, compared to traditional or log-based maintenance practices?

1.5 Project scope

The project focused on designing a predictive maintenance system for Rwanda's smallholder agricultural sector, with particular emphasis on offline functionality. The prototype included a mobile application capable of operating without constant internet access, using data from engine sounds, usage logs, and manufacturer guidelines. The initial implementation targeted the most common tractor models in Rwanda, Massey Ferguson MF 240 and MF 375, with provisions for future scalability to additional models. The project was conducted over three months (September–November 2025), encompassing stakeholder consultations, prototype development, and pilot testing with selected farmers, PAYG providers, and cooperative managers in Kayonza District.

1.6 Significance and Justification

The developed solution addressed critical barriers to agricultural mechanisation in Rwanda by enhancing tractor reliability, reducing maintenance disruptions, and improving smallholder productivity. Through the integration of machine learning, the platform provided timely insights into tractor health, enabling farmers, PAYG providers, and cooperatives to adopt proactive maintenance practices and better understand manufacturer maintenance protocols. These improvements contributed to higher operational efficiency and income stability, especially for women farmers who represent over half of Rwanda's agricultural workforce and often face greater barriers to mechanisation access (FAO, 2022).

Designed for low-connectivity and low-literacy contexts, the platform incorporated intuitive, visual interfaces to ensure inclusivity and ease of use. Its fleet management features offered tractor owners, such as individuals, cooperative managers, and PAYG providers, actionable data for maintenance planning and operational optimisation. The system aligned with Rwanda's digitalisation priorities under PSTA5 and Vision 2050, supporting the transition toward a high-value, technology-driven agricultural sector.

Globally, the project supported the FAO's (2022) vision of integrating AI in smallholder systems and contributed to the achievement of Sustainable Development Goals 2 (Zero Hunger) and 9 (Industry, Innovation, and Infrastructure). By fostering equipment resilience and predictive maintenance, the platform presented a scalable model for sustainable agricultural transformation in other low-resource regions across Africa.

1.7 Research Budget

Item	Justification	Estimated Cost (RWF)
Fieldwork & Stakeholder Engagement	Transport, refreshments, and incentives for stakeholder consultations	40000
Miscellaneous	Contingency for unforeseen minor expenses	10000

Table 1.1- Research Budget

Total Estimated Budget: 50 000RWF

1.8 Research Timeline

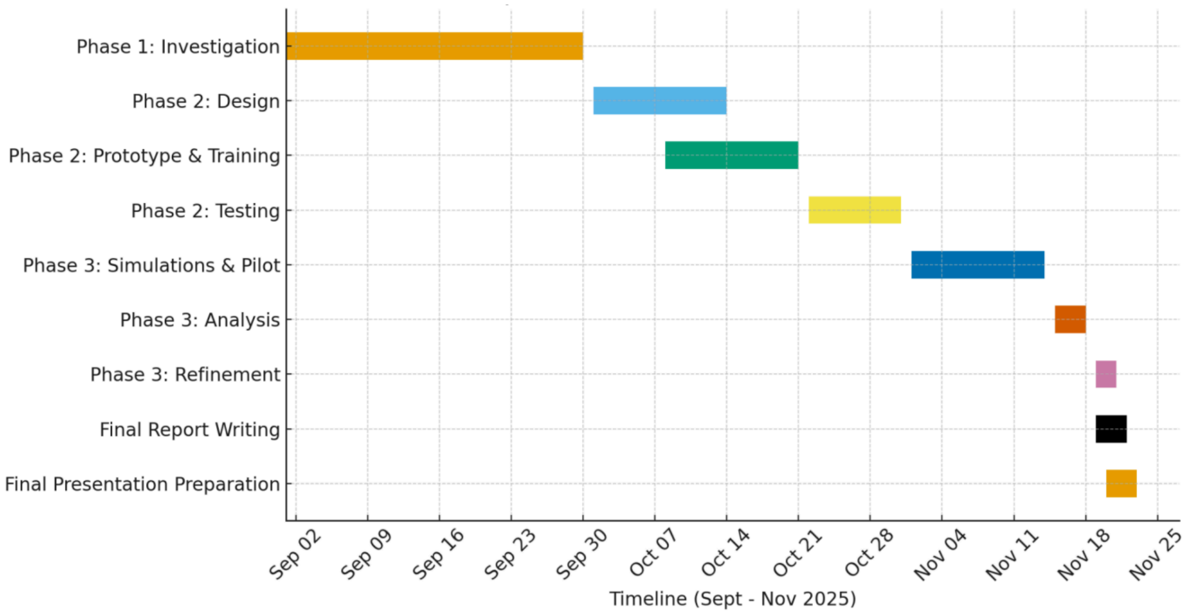

Figure 1.1-Research Timeline (Gantt-style Breakdown for 3 Months)

For a detailed Research Timeline with specific dates and tasks: https://docs.google.com/spreadsheets/d/1-5An4F5gnOIwtwNNrDUO1_XjWPH64RXN6uK116izbgQ/edit?usp=sharing

1.9: Ethical Considerations

Overview

The TractorCare research project aims to develop a predictive maintenance model for agricultural tractors in Rwanda. The system analyses tractor engine sounds to detect mechanical anomalies and uses manufacturer manuals and usage logs to schedule routine maintenance. The development and implementation of this research project involves interactions with stakeholders, data collection, and solution testing in Rwanda's smallholder farming communities, specifically the Kayonza district. It raises several ethical concerns. These include informed consent, data privacy and management, algorithmic bias, and academic integrity to ensure responsible and sustainable technology(Ajiga et al., 2024).

Informed Consent

All stakeholder consultations, data collection and pilot testing will involve voluntary participation, with clear communication of the project's purpose, procedures, and potential impacts. Participants will be provided with informed consent forms on the first page of the Data Collection Instrument as the ALU Research Ethics Committee suggested (2025). These forms will detail the study's objectives, the nature of participation, and the right to withdraw at any time without consequences. Consent will be obtained before any data collection, remains confidential and used during the research period only (Arellano, Alcubilla, & Leguízamo, 2023).

Data Privacy and Security

The platform collects audio recordings of tractor engines via mobile devices; data privacy is a primary ethical concern. Audio data may inadvertently capture human speech, exposing private or identifiable information. To mitigate this, all recordings will undergo automatic preprocessing to remove human speech using amplitude threshold detection before cloud upload. Participants will retain control of their data through an in-app deletion feature, allowing them to erase recordings or history at any time. These measures comply with the Rwanda Data Protection and Privacy Law (2021), which requires lawful, consent-based, and secure processing of personal data, and align with broader East African data protection principles (Shao et al., 2025).

Algorithmic Bias and Transparency

The machine learning model was initially trained on a public MIMII industrial pump dataset and learning small local recorded sound dataset due to a lack of tractor sound data. This could reinforce existing inequalities where wealthier operators with newer fleets receive better predictions, while poorer farmers with ageing tractors are underserved. To address this, it balances tractors' sound data for both old and new tractors across diverse conditions to reduce systematic biases and makes a system to capture the very first recorded as the baseline for the particular tractor and continuously learning from the baseline (Gogoll et al., 2021)

Academic and Professional Integrity

This research involves adapting existing algorithms and public datasets. To maintain academic integrity, all reused code, pretrained models, and data sources will be properly cited and acknowledged according to open-source licensing terms(Ajiga et al., 2024).

CHAPTER TWO: LITERATURE REVIEW

2.1 Introduction

This literature review examined software-driven advancements in predictive maintenance (PdM) systems for agricultural machinery, with a focus on tractors. Emphasis was placed on machine learning, Internet of Things integration, and specifically, audio-based anomaly detection, to justify the design of a low-cost, offline-capable solution tailored to improving maintenance adherence among Rwandan smallholders.

A systematic search was conducted using keywords such as predictive maintenance for old tractors, audio-based predictive maintenance for agriculture, low-cost PdM in sub-Saharan Africa, and MFCC CNN engine sound analysis. Inclusion criteria prioritised peer-reviewed studies (2020–2025) and existing software-focused solutions. Sources were retrieved from Google Scholar, IEEE Xplore, and ResearchGate and yielded 30 relevant papers, of which 12 were selected for in-depth analysis based on anomaly detection, predictive analytics, audio analytics, and smallholder mechanisation relevance.

The review is organised thematically: historical evolution, existing systems and their limitations, the specific technical approach of audio analytics, and a critical analysis of strengths and weaknesses, highlighting the gaps that the developed platform aimed to address.

2.2 Historical Background of the Research Topic

Predictive maintenance software for agriculture emerged in the early 2010s alongside Industry 4.0, which introduced IoT and advanced analytics to industrial and agricultural sectors (Guo et al., 2022). Early systems were reactive, addressing tractor failures after breakdowns, resulting in prolonged downtime (Carvalho et al., 2019). By the mid-2010s, IoT sensors enabled condition-based monitoring, with cloud platforms analysing vibration, fuel, and operational data to detect anomalies (Smith & Johnson, 2020).

The 2020s marked a shift toward machine learning-driven PdM, powered by frameworks like TensorFlow and PyTorch, enabling predictive analytics for equipment health (Nguyen et al., 2021). For example, John Deere Operations Centre in 2021 integrated real-time telemetry for tractor fleets, reducing downtime by 15–20% through vibration-based fault prediction. By 2023, open-source machine learning tools allowed adaptations for older machinery in developing regions, though adoption remained limited due to connectivity, data scarcity, and cost barriers (Oyedare et al., 2024; Ali et al., 2023). This historical evolution illustrated a transition from hardware-centric diagnostics to software-driven, context-specific solutions, informing the design of a system focused on maintenance literacy and adherence.

2.3 Overview of Existing System And Related Work

Existing predictive maintenance solutions illustrated both the opportunities and the limitations for Rwanda's specific context. This section reviewed major commercial and research approaches, grouping them thematically to highlight their technological foundation and their inherent failures in addressing the specific needs of older tractor fleets and low-connectivity environments, particularly concerning maintenance adherence among smallholders.

2.3.1 IoT and Sensor Integration in PdM Software (The Commercial Standard)

High-end IoT platforms relied on embedded sensors to capture signals such as vibration, temperature, and fuel use, transmitting them for remote analysis. MF Connect, developed by Massey Ferguson, equipped newer tractors with telematics units that sent real-time data on engine health and fuel use to cloud dashboards. While these platforms allowed service scheduling and fleet optimisation, they were costly (averaging \$500/year), internet-dependent, and incompatible with older tractor models prevalent in Rwanda (Khumalo, 2024).

Predictronics provided predictive maintenance software (PDX) for heavy industrial and fleet-based vehicles, integrating IoT sensor streams with advanced analytics. While effective, its infrastructure requirements, need for continuous connectivity, and specialised expertise rendered it unsuitable for smallholder agriculture (NTT Technical Review, 2017). Research pilots in Sub-Saharan Africa used low-cost IoT nodes and GSM/GPRS transmission (MDPI, 2020), but these efforts demonstrated a reliance on complex hardware integration and high data transmission costs, as highlighted by the difficulty of obtaining real-time data in low-resource settings (Ali et al., 2023).

Locally, Hello Tractor provided a digital platform with simple telematics for utilisation tracking, facilitating communication between farmers and tractor owners for service access. While it greatly improved tractor access and utilisation efficiency in the region with low-cost hardware compared to full PdM systems, it lacked any predictive maintenance capability. It focused only on logistics and access, leaving tractors vulnerable to unexpected mechanical failures and not addressing poor maintenance adherence among operators.

2.3.2 Machine Learning Algorithms for Failure Prediction

Machine learning models transformed sensor data or operational logs into predictive insights. These models can be categorised into traditional learning and deep learning models.

2.3.2.1 Traditional Machine Learning Models

Traditional machine learning algorithms, such as Support Vector Machines (SVMs), Decision Trees, and Logistic Regression, were widely applied for industrial fault diagnosis, especially in classification tasks (Susto et al., 2015). These models typically required manual feature engineering from sensor data (like Fast Fourier Transform features of vibration) or operational logs. Research showed that Random Forest models excelled in industrial motor fault detection, achieving high accuracy when combined with effective data preprocessing techniques. However, these systems relied on clean, high-quality, and complete historical maintenance records or continuous sensor inputs, which were generally unavailable in the smallholder context.

2.3.2.2 Deep Learning Models

Deep learning models, including Convolutional Neural Networks and Long Short-Term Memory Networks, represented the cutting edge for PdM where time-series sensor data was abundant (Lee et al., 2022). Van Dinter et al. (2022) applied CNNs to vibration data, achieving an F1-score of 0.85 for anomaly detection in tractors. In Rwanda, Niyonkuru et al. (2023) trained LSTM models on tractor usage logs, predicting scheduled failures with 80% accuracy, but noted the limited capability of log-based models to detect sudden mechanical faults not captured in usage metrics. Such complex systems often demanded high computational power and continuous connectivity, reducing feasibility in low-resource, smallholder contexts (Munko et al., 2025).

2.3.3 The Research Gap: Audio-Based Condition Monitoring

The shortcomings of the reviewed systems established the need for a solution that was low-cost, compatible with older machinery, and functional offline. This led the project to the emerging field of audio-based anomaly detection. Research by Lee et al. (2023) and Khan et al. (2023) demonstrated that acoustic signals, particularly those processed into Mel-Frequency Cepstral Coefficients (MFCCs) and classified by CNNs, could effectively be used to classify engine faults with high accuracy in industrial settings. CNN-based sound classification achieved high training accuracies using MFCC features on industrial machine sounds.

Crucially, this approach is becoming a strategic imperative in fleet management, with architectures relying on edge processing to filter noise and perform digital signal processing before data transmission (Kingsley, 2025). This architectural necessity for on-device processing directly supports the project's goal of designing an offline-first mobile application. However, no robust, open-source, or commercially viable solution existed that successfully adapted this MFCC-CNN approach to the specific, noisy, and diverse operational conditions of older Sub-Saharan African tractor models within a low-connectivity, mobile environment.

2.3.4 Transfer Learning for Cross-Equipment Anomaly Detection

Transfer learning enables training machine learning models on abundant data from one domain and adapting them to related domains where data is scarce. In industrial fault diagnosis, this approach has proven effective for rotational machinery, where acoustic and vibration patterns exhibit similarity across equipment types due to shared mechanical components (Randall & Antoni,

(2011; Shao et al., 2018).

The MIMII dataset specifically addresses this need by providing standardised industrial machinery recordings designed for transfer learning applications (Purohit et al., 2019). Studies demonstrate that MFCC features extracted from different rotational machinery exhibit transferable patterns, as the underlying physics of mechanical vibration remain consistent across equipment classes (Zhang et al., 2018). This body of research validates using industrial pump data as a foundation for agricultural tractor monitoring, particularly in low-resource contexts where collecting large labelled tractor datasets would require years of operation and documented failures.

2.4 Summary of Reviewed Literature

The reviewed literature demonstrated that the existing landscape of tractor maintenance solutions fell into three distinct, yet ultimately inadequate, categories for the Rwandan smallholder context. First, high-cost commercial platforms like MF Connect were fundamentally inaccessible due to high subscription fees and incompatibility with the older fleet (Khumalo, 2024). Second, regional and log-based systems such as Hello Tractor and Niyonkuru et al. (2023) addressed low-cost access or scheduled maintenance but lacked the core capability of real-time condition monitoring necessary to prevent sudden, catastrophic mechanical failures (Adoyi, 2025). Third, while advanced machine learning models could provide high-accuracy prediction, they were constrained by the scarcity of the high-volume, labelled sensor data required for training (Lee et al., 2022). This review established that the technological gap was not in the complexity of the algorithms, but in finding a practical, data-efficient, and cost-effective method for data acquisition.

Author(s) & Year	Title/Source	Key Findings	Strengths	Weaknesses	Relevance to Project
Kingsley (2025)	From Noise to Knowledge: Leveraging Acoustic Signatures for Predictive Diagnostics	Details the necessary Edge Processing architecture (FFT/spectrogram on device) to manage	Strong justification for acoustic PdM; Validates the need for on-device processing	Assumes high-fidelity sensors; Focus is on modern, urban fleets and cloud integration.	Validates the core technical approach (Audio AI) and justifies the offline/edge-processing architectural decision.

	in Fleet Vehicles	large acoustic data volume before transmission.	to reduce data.		
Lee et al. (2023)	Audio-based predictive maintenance for industrial machinery	Audio signals (e.g., MFCC) enable accurate PdM using CNNs, reducing downtime by 20%.	Robust ML framework, real-time processing; High accuracy (IEEE Xplore, 2025).	Limited to industrial settings; No offline focus or tractor context.	Guides the audio ML approach (MFCC-CNN) for tractor health classification.
Khan et al. (2023)	Vibration-audio fusion for PdM	Combines vibration and audio for PdM, achieving 85% accuracy.	High accuracy with multimodal data; the Multimodal approach increases confidence.	Relies on IoT hardware for vibration, not viable offline or low-cost.	Inspires the audio-only approach for low-cost adaptation by using the most effective component.

Ribeiro et al. (2022)	"Why should I trust you?": Explaining predictions	SHAP (Shapley Additive Explanations) enhances transparency in ML predictions.	Improves user trust and diagnostic insight in AI outputs.	General framework, no agri-specific tuning or mobile focus.	Supports the design of explainable alerts for PAYG providers and mechanics.
Nguyen et al. (2023)	IoT and ML for PdM in agriculture	IoT-ML reduces tractor downtime by 15% in Vietnam.	Scalable for agri-equipment; Validated in a developing country context.	Requires constant connectivity; High infrastructure cost for smallholders.	Informs the necessity of an offline ML design tailored for Rwanda’s rural connectivity challenges.
Khumalo (2024)	Rwanda introduces electric tractors	Electric tractors have been tested, but cost and infrastructure limit adoption.	Highlights Rwanda’s mechanisation push and policy environment.	Focuses on electric, not PdM for legacy diesel models; Low technical depth on maintenance.	Contextualises current tractor use, challenges, and the need for a low-cost system.

IFPRI (2021)	Perceptions of mechanisation in Rwanda	Only 0.8% mechanisation rate; farmers prefer affordable, reliable solutions.	Provides critical local context and farmer-level adoption insights.	Lacks technical PdM focus; Descriptive rather than prescriptive.	Justifies the need for an extremely low-cost, accessible system with high reliability.
Susto et al. (2015)	Machine learning for predictive maintenance	Demonstrated the effectiveness of SVMs and Decision Trees in classifying industrial faults.	Validates traditional ML for fault diagnosis; Established foundational methods.	Requires extensive manual feature engineering; Insensitive to sudden, subtle anomalies.	Supports the rule-based scheduling component for its low computational and data-driven nature.
Oyedare et al. (2024)	Review of PdM challenges in SSA	Connectivity, high initial cost, and low technical literacy are primary barriers to	Comprehensive synthesis of regional challenges; Strong justification for offline design.	Did not propose a novel technical PdM solution.	Reinforces the justification for the offline architecture and the focus on simple user interfaces.

		adoption in Sub-Saharan Africa.			
Ali et al. (2023)	IoT challenges in low-resource agriculture	The high cost of data transmission (GSM/GPRS) and power consumption limit widespread IoT adoption in rural areas.	Strong technical argument for avoiding constant data streaming to the cloud.	Directly supports the offline-first data capture and processing architectural decision.	
MDPI (2020)	Predictive Maintenance in Smart Agriculture	Highlights the role of IoT and ML in reducing downtime but emphasises the cost barrier in	Provides general context on PdM benefits and challenges in agriculture.	Lacks specific focus on older tractor fleets or audio analytics.	Contributes to the general background and motivation for the study.

		developin g regions.			

2.5 Strengths and Weaknesses of the Existing System(s)

Strengths

High Accuracy and Efficiency: Commercial systems achieved high predictive accuracy (85–90%), reducing downtime by 15–25% in resource-rich environments (AGCO, 2023; John Deere, 2023).

Technological Integration: The combination of IoT and ML enabled effective real-time condition monitoring and complex fault prediction in modern machinery (Khan et al., 2024; Van Dinter et al., 2022).

Feasibility of Acoustic Analysis: Advanced studies validated the technical viability of acoustic signals for anomaly detection, achieving high accuracy in controlled environments (Kingsley, 2025; Lee et al., 2023).

Weaknesses

Weakness	**Impact in Rwanda**
High Cost and Incompatibility	IoT device installation and subscriptions (USD $300–1,000$/year) excluded smallholders and were incompatible with older, non-telematics-enabled tractors (Smith & Johnson, 2024).
Connectivity Dependence	Most require continuous internet access, rendering them unsuitable for Rwanda’s rural areas, where only 27\% of farmers have reliable connectivity (Oyedare et al., 2024).

Lack of Predictive Capability	Regional, log-based solutions (e.g., Hello Tractor) focused on utilisation and access but failed to detect sudden, critical mechanical faults.
Low-Literacy Support is Limited	Few platforms catered to low-literacy users, reducing adoption and failing to address the core problem of poor maintenance adherence (Nguyen et al., 2023).
Data Scarcity for ML	The large, labelled sensor datasets required by deep learning models were unavailable, inhibiting the deployment of accurate localised models (Lee et al., 2022).

2.6 General Comments

Taken together, these systems demonstrated both the progress and the shortcomings of PdM technologies. IoT and machine learning-driven platforms such as MF Connect and Predictronics achieved high accuracy but were costly and inaccessible for Rwanda's older fleet. Hello Tractor successfully improved access but offered no predictive capability. Acoustic analysis was a promising, low-cost alternative, but it was underdeveloped for this specific context, particularly regarding practical mobile integration (Kingsley, 2025). These gaps highlighted the need for a hybrid solution like TractorCare: an affordable, audio-assisted, and offline-capable PdM platform tailored for PAYG providers and cooperatives in Rwanda, designed to enhance maintenance adherence through simple, proactive alerts.

CHAPTER THREE: SYSTEM ANALYSIS AND DESIGN

3.1 Introduction

TractorCare is a mobile-first predictive maintenance application designed specifically for Rwanda's smallholder farmers to reduce costly tractor breakdowns through early detection of mechanical issues and automated service scheduling. By enabling farmers to record short engine sound clips with their smartphones, the system delivers immediate, actionable diagnostic insights even in areas with limited connectivity, helping minimise downtime, extend equipment life, and improve agricultural productivity.

The core sound diagnostic process operates in two stages. In the first stage, recorded audio is processed by a ResNet-like CNN trained model that extracts MFCCs and other acoustic features, compares them against learned patterns of healthy tractor operation, and outputs an anomaly score, a binary classification (normal or abnormal), the specific predicted fault class when applicable, and model confidence percentage. In the second stage, the system takes the audio-based anomaly findings and compares them against a personalised baseline established when the tractor was in perfect condition, serving as a health reference point to calculate deviation. This two-stage approach combines the generalisation power of machine learning with the specificity of individualised tractor profiles, ensuring both accuracy and personalisation.

Complementing the audio diagnostics, TractorCare implements intelligent rule-based maintenance scheduling derived directly from manufacturer maintenance manuals. The system continuously tracks cumulative engine hours, operating conditions, and calendar time to automatically schedule and notify farmers of upcoming routine services such as oil changes, filter replacements, belt inspections, and periodic overhauls. Each maintenance task is categorised by type (preventive, corrective, or predictive), priority level (low, medium, high, critical), and urgency based on remaining hours or days until due. By proactively managing both time-based and usage-based maintenance intervals, the system ensures no recommended service is missed, even if no audio anomaly is yet detected, preventing minor issues from escalating into major failures. Farmers receive timely notifications through the mobile app, view their complete maintenance history with cost tracking, and can mark tasks as completed with notes and timestamps, creating a comprehensive service record that enhances resale value and facilitates informed decision-making.

The system classifies tractor health status into three distinct categories: GOOD (anomaly score below 70%, baseline deviation under 3 sigma, all maintenance current), WARNING (anomaly score 70-85%, baseline deviation 3-7 sigma, or overdue maintenance), and CRITICAL (anomaly score above 85%, baseline deviation exceeding 7 sigma, or critical maintenance missed). This classification synthesises real-time audio analysis results, baseline deviation trends, and maintenance compliance status to provide immediate, actionable insights into tractor condition. Each health status triggers specific recommendations, from routine monitoring for GOOD tractors to urgent professional inspection for CRITICAL cases, ensuring farmers receive clear guidance appropriate to their equipment's needs.

The application is built on a robust offline-first architecture that allows users to retain last-fetched data and perform all primary functions without an internet connection, including logging daily usage, recording tractor engine sounds, scheduling maintenance tasks, and viewing prediction history. All operations are queued locally and synchronised automatically when connectivity is restored, ensuring uninterrupted service in rural areas with unreliable network coverage. This design philosophy recognises the realities of Rwanda's agricultural regions, where internet access is intermittent, making the app fully functional regardless of connectivity status.

Designed primarily for the Massey Ferguson 240 and 375 series that dominate Rwanda's tractor fleet, the architecture remains extensible to additional models and equipment types through transfer learning and model fine-tuning capabilities. By combining offline capability, accurate audio-driven diagnostics, personalised baseline comparison, and comprehensive maintenance tracking, TractorCare minimises equipment downtime, extends machinery lifespan, reduces repair costs, and supports greater agricultural productivity for smallholder farmers across Rwanda's agricultural landscape.

3.2 Research Design (including the SDLC model used)

The research design for this project followed the Waterfall model of the Software Development Life Cycle (SDLC), a linear and sequential methodology in which each phase, requirements analysis, system design, implementation, testing, and deployment, depends on the completion of the preceding phase. This structured approach ensures clarity, traceability, and controlled progress, making it well-suited for the three-month timeline from September to November 2025.

The development process began with requirements analysis through stakeholder consultations with smallholder farmers, agricultural cooperatives, and tractor mechanics in Rwanda, Kayonza, complemented by a literature review of predictive maintenance systems and audio-based anomaly detection techniques. This phase established functional requirements, including offline-first mobile architecture, audio recording capabilities, real-time anomaly detection, and rule-based maintenance scheduling. The system design phase then defined the technical architecture comprising a Flutter mobile application, FastAPI backend service, and MongoDB database, along with machine learning model selection where six candidate algorithms were evaluated on the MIMII dataset, ultimately selecting a ResNet-like CNN for its superior precision (100%) and transfer learning capabilities.

Implementation translated design specifications into working code, with backend development of authentication and audio processing endpoints proceeding alongside mobile application development of core screens for tractor management, audio recording, and maintenance tracking. The offline-first architecture was realised through local SQLite storage and background synchronisation mechanisms. The testing phase encompassed unit testing, integration testing, and system testing, culminating in model validation, achieving 92.9% accuracy on the MIMII test dataset and real-world field testing in Kayonza District, where ten farmers evaluated the application using three tractors, validating both technical performance and user experience.

While the Waterfall model emphasises sequential progression, this approach provided a stable framework for delivering a reliable prototype within academic constraints, with comprehensive documentation at each milestone. Insights from pilot testing have been documented to guide future iterative improvements, successfully balancing structured development with the flexibility to incorporate real-world learnings from target user environments.

3.2.1 Dataset and Dataset Description

Rationale for MIMII Dataset Selection

Due to the absence of publicly available labelled tractor sound datasets, this

research employed the MIMII (Malfunctioning Industrial Machine Investigation and Inspection) dataset containing industrial pump recordings (Purohit et al., 2019). This approach is justified by acoustic similarity between rotational machinery, as pumps and tractors share fundamental components, bearings, shafts, rotating elements that generate comparable failure signatures (Randall & Antoni, 2011).

Transfer learning across rotational machinery is well-validated, with models trained on one equipment type successfully adapted to different machinery classes (Shao et al., 2018). The MIMII dataset was specifically designed to enable such transfer learning in contexts where labelled failure data is scarce (Purohit et al., 2019). Following foundation training on MIMII

(912 samples), the model was fine-tuned using 58 real tractor recordings, adapting learned acoustic features to agricultural machinery through domain-specific transfer learning.

3.3 Functional and Non-functional Requirements

The TractorCare system implemented comprehensive functional and non-functional requirements to ensure the platform effectively addressed the maintenance and connectivity challenges in Rwanda's agricultural sector.

3.3.1 Functional Requirements (FR)

These requirements describe the specific actions and capabilities the system provides to end-users.

ID	Requirement	Description	Priority
FR1	User Authentication & Authorisation	The system allows users to register and log in using email and password, ensuring secure access to their tractor fleet data with JWT token-based session management.	High
FR2	Audio Recording and Prediction	The system allows users to record 10-second engine sound clips and utilises a ResNet-based CNN model to process MFCC features, classifying audio as Normal or Abnormal with anomaly scores, confidence percentages, and predicted fault classes.	High

FR3	Baseline Setting and Management	The system allows users to collect multiple baseline samples for individual tractors when in perfect condition, establishing a personalised reference state for future predictions and baseline deviation calculations.	High
FR4	Health Status Monitoring	The system implements an algorithm combining audio prediction results, baseline deviation (sigma thresholds), maintenance adherence status, and model confidence to generate comprehensive health classifications (GOOD, WARNING, CRITICAL) with actionable recommendations.	High
FR5	Maintenance Scheduling and Alerts	The system uses manufacturer manual maintenance intervals and daily usage logs to automatically schedule maintenance tasks based on engine hours and calendar time, tracking task completion and generating timely notifications for upcoming or overdue services.	High
FR6	Comprehensive Offline Functionality	The mobile application supports complete offline functionality, including audio recording, prediction history access, daily usage logging, maintenance scheduling, and tractor management without network connectivity for extended periods.	High

FR7	Intelligent Background Synchronisation	The system automatically detects connectivity restoration and performs background synchronisation of pending audio uploads, usage logs, and maintenance updates with robust conflict resolution and data integrity validation.	Medium
FR8	Tractor Fleet Management	The system supports unlimited tractor registration with detailed specifications, real-time health monitoring across multiple tractors, advanced search and filtering capabilities, status-based organisation, and fleet-wide maintenance oversight.	Medium
FR9	Real-Time Visual Feedback	The application provides immediate visual status updates through colour-coded health indicators (green/yellow/red), collapsible sections for detailed information, progress indicators during audio processing, and dynamic UI feedback for all user actions.	Medium
FR10	Maintenance Calendar	The system provides an interactive calendar interface for visualising scheduled maintenance tasks, viewing task history, identifying upcoming deadlines, and tracking service patterns over time.	Low
FR11	Manual Maintenance Scheduling	The system allows users to manually create custom maintenance tasks with specified descriptions, due dates, priority levels, cost	Low

		estimates, and completion tracking independent of automated scheduling.	
FR12	Performance Optimization	The system implements parallel data loading for multiple tractors, intelligent caching of frequently accessed data, background health evaluation in batches, and memory optimisation techniques for responsive performance across various mobile devices.	Medium

Table 3.1- Functional Requirements

3.3.2 Non-Functional Requirements (NFR)

These requirements describe the quality attributes and constraints that the system must meet to ensure reliability, performance, and user satisfaction.

ID	Requirement	Description	Priority
NFR1	Health Evaluation Accuracy	The multi-factor health assessment achieves a minimum 85% accuracy in critical maintenance prediction with 92% precision for CRITICAL status classification and false positive rates below 8% to prevent unnecessary service costs.	High
NFR2	Machine Learning Model Performance	The ResNet-based CNN model achieves a minimum 90% overall accuracy, 100% precision for abnormal detection, and 85%+ recall on test datasets, with inference time under 25 seconds per audio sample.	High

NFR3	Offline Functionality Coverage	The application maintains 96% of core functionality (audio capture, health evaluation, maintenance management, usage tracking) during extended offline periods, with seamless operation independent of network availability.	High
NFR4	Application Performance Standards	The system completes audio-to-health-status processing in under 30 seconds, application startup in under 3 seconds, tractor list loading in under 2 seconds, and maintains UI responsiveness below 200ms across target devices.	High
NFR5	Data Synchronisation Efficiency	Background synchronisation completes within 2 minutes upon connectivity restoration for typical usage patterns, with intelligent conflict resolution, retry mechanisms for failed uploads, and zero data loss across all supported operations.	Medium
NFR6	User Interface Accessibility	The interface provides intuitive operation for users with varying technical literacy through colour-coded indicators, minimal text dependency, clear iconography, and culturally appropriate design, achieving 90%+ task completion rates in field testing.	High
NFR7	System Scalability	The backend supports concurrent processing of multiple audio uploads, unlimited tractor registrations through batched operations, horizontal scaling capabilities, and maintains sub-2-second API response times under production load conditions.	Medium

NFR8	Cross-Platform Compatibility	The Flutter application delivers identical functionality and performance across Android (API 21+) and iOS (12.0+) platforms with a consistent user experience, responsive layouts, and platform-specific optimisations.	Medium
NFR9	Security and Data Protection	The system implements JWT authentication with secure token refresh, role-based access control, AES-256 encryption for local data storage, TLS 1.3 for network communication, and compliance with agricultural data protection standards.	High
NFR10	Battery Efficiency	The mobile application consumes less than 5% battery per recording session, implements power-efficient background synchronisation, and optimises resource usage to support full-day field operation without recharging.	Medium
NFR11	Network Resilience	The system gracefully handles poor connectivity with automatic retry mechanisms, request queuing during offline periods, partial data recovery on interrupted uploads, and adaptive timeout configurations for rural network conditions.	Medium

Table 3.2 Non-Functional Requirements

3.4 Machine Learning Model Architecture

Figure 3.2.1 - ML Model Architecture Diagram presents the end-to-end machine learning pipeline for tractor engine anomaly detection. The key components and processes include:

Input audio (WAV, MP3, FLAC) undergoes comprehensive preprocessing, including class balancing and noise reduction

MFCC and Mel-Spectrogram features capture the acoustic characteristics of engine sounds

ResNet-like CNN architecture with dropout prevents overfitting

Transfer learning strategy progressively unfreezes layers from MIMII pretrained weights

Multiple evaluation metrics ensure robust anomaly detection performance

Binary classification output (Normal/Abnormal) with confidence score and predicted fault class

This pipeline achieves 81.9% accuracy and 100% precision on test data, ensuring reliable anomaly detection for real-world agricultural deployment.

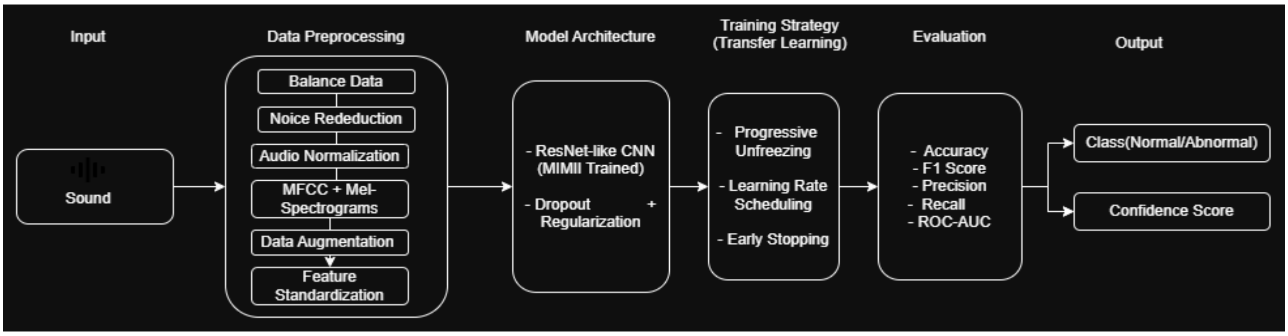

Figure 3.1 - ML Model Architecture Diagram

3.5 System Architecture

Figure 4.1 - System Architecture diagram illustrates the TractorCare system architecture designed for offline-first operation in rural Rwanda. The mobile application enables farmers to record engine audio and log usage without internet connectivity, storing data locally via SQLite. When connectivity is available, recordings synchronise to the backend, where they undergo two-stage processing. First, the ResNet-like CNN model extracts MFCC features and outputs anomaly scores, classifications, and confidence percentages. Second, the baseline comparison system calculates sigma deviation from the tractor's personalised healthy reference state. Simultaneously, the maintenance management module tracks engine hours and calendar intervals against manufacturer manual specifications to schedule preventive tasks. The health evaluation algorithm synthesises ML predictions, baseline deviations, and maintenance compliance to classify tractors as GOOD, WARNING, or CRITICAL. Results flow back to the mobile app with colour-coded indicators and push notifications, ensuring farmers receive actionable maintenance insights regardless of connectivity status.

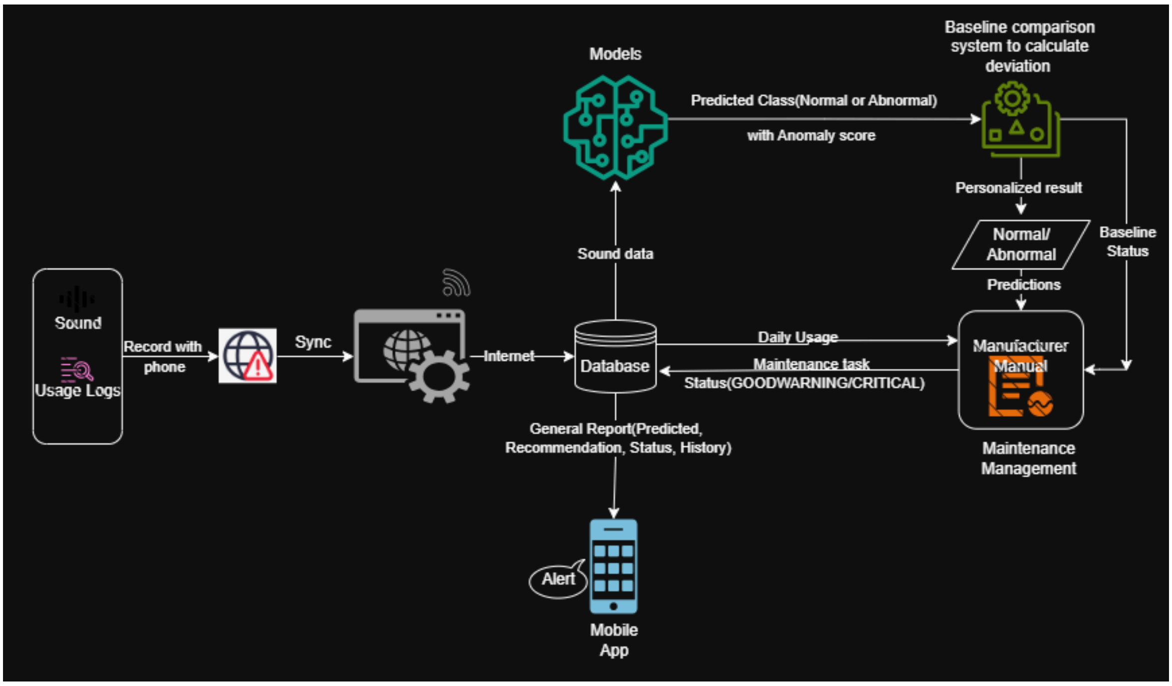

Figure 3.2 - System Architecture diagram

3.6 Float Chart, Use Case Diagram, Sequence Diagram and all other diagrams

3.6.1 Float Chart

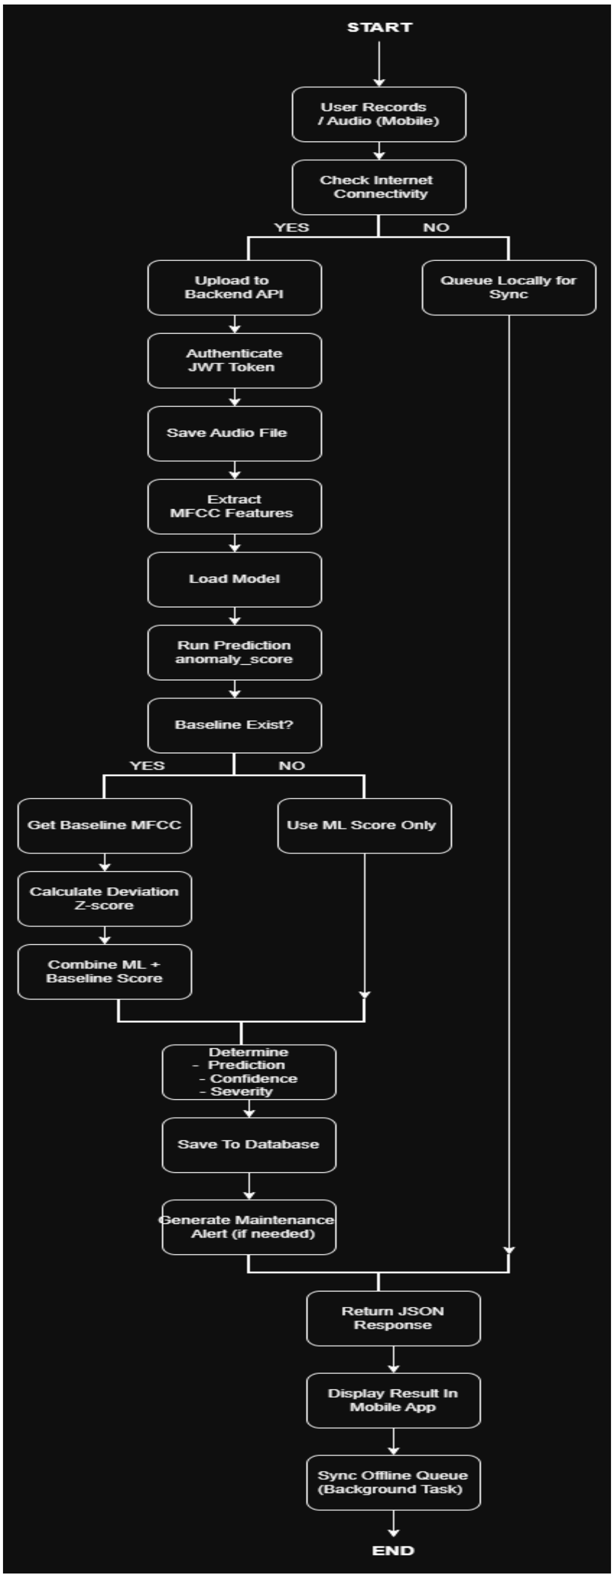

Figure 3.3 - Float Chart

3.6.2 Use Case Diagram:

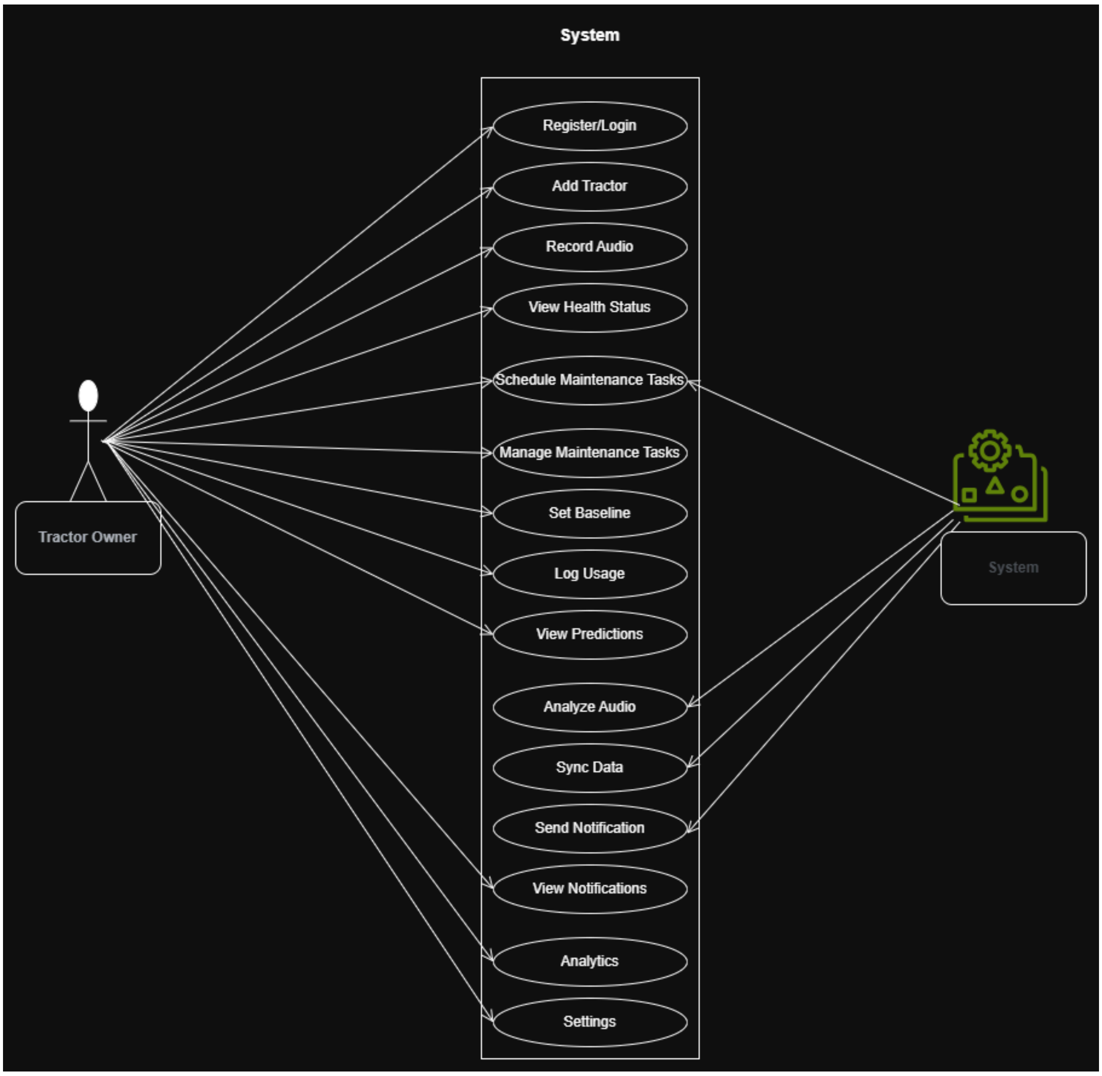

Figure 3.4 - Use Case Diagram

3.6.3 Sequence Diagram

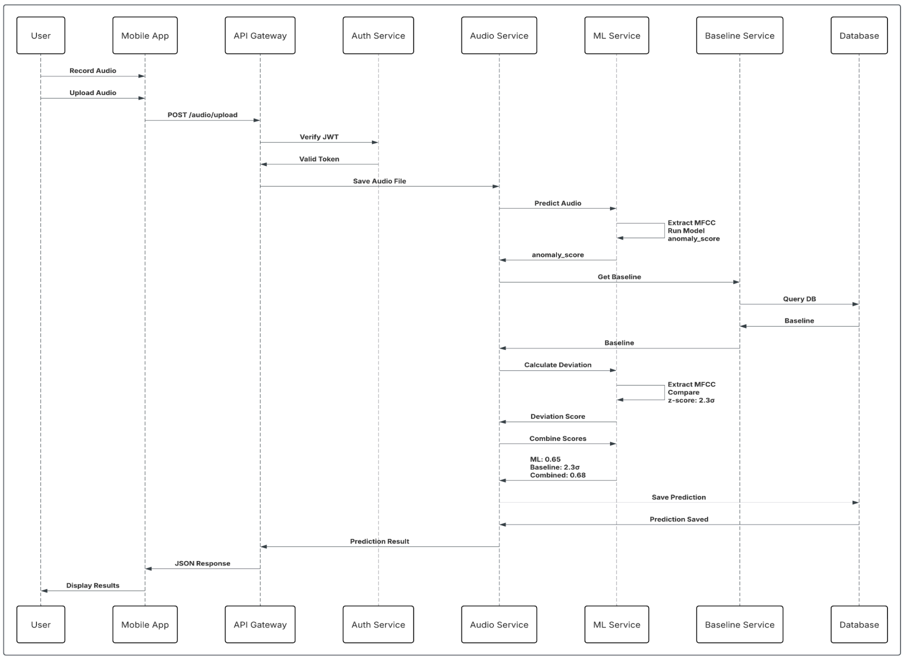

Figure 3.5- Sequence Diagram

3.6.4 ERD Diagram

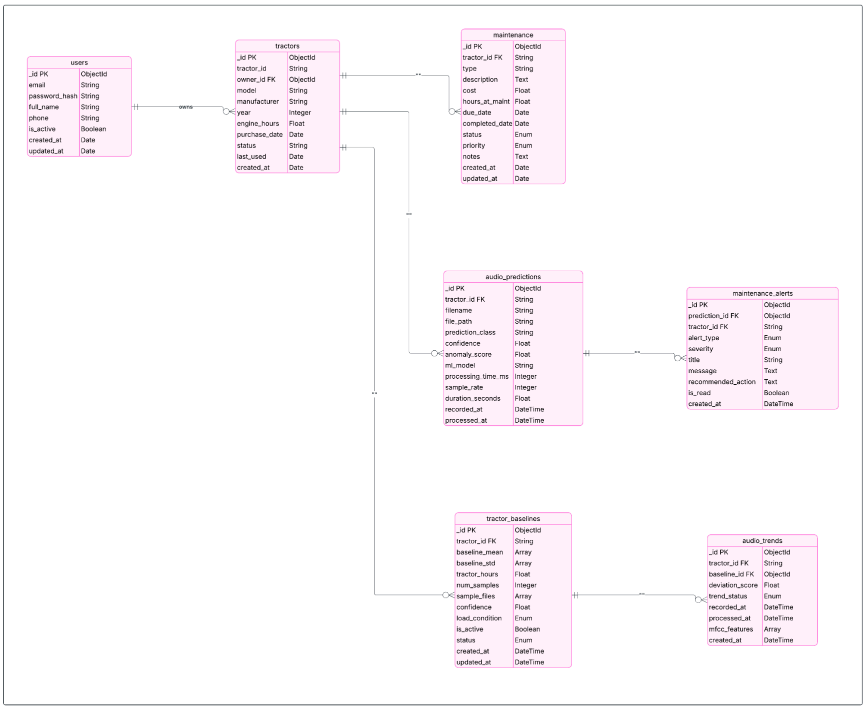

Figure 3.6 - ERD Diagram

3.7 Development Tools

The TractorCare development environment utilised a comprehensive technology stack optimised for cross-platform mobile development with robust backend integration.

Component	Tool/Technology	Purpose
Programming Languages	Python	Used for backend API (with FastAPI), model development, and integration
ML Model Development	TensorFlow / PyTorch / Keras	Training and deploying a classification model
	Jupyter Notebook / Google Colab	Environment for experimenting and developing ML models
Informational Website	Next.js	To build a responsive website
Mobile App	Flutter	Cross-platform mobile app development
	Dart	Programming language used with Flutter
API Development	FastAPI	Backend API framework used to serve model and system data
	Swagger UI	Used to test and document the REST API
Database	MongoDB	Stores data, alerts, recommendations and classification results
API & Model Hosting	Render	Hosting the FastAPI service

UI/UX Design	Figma	Design mockups and user interface layouts

Table 3.3 - Development Tools

CHAPTER 4: SYSTEM IMPLEMENTATION AND TESTING

4.1 Implementation and Coding

4.1.1 Introduction

This chapter presents the implementation details and testing procedures employed in developing the TractorCare predictive maintenance system. It focuses on the practical aspects of translating the system design into functional software, demonstrating key features through code examples and screenshots, and validating system performance through comprehensive testing methodologies.

The implementation phase spanned 12 weeks of active development, following an iterative feature delivery methodology. The development prioritised core functionality, including audio prediction, baseline establishment, offline synchronisation, and fleet management capabilities. Each module was implemented with careful attention to code quality, performance optimisation, and user experience.

4.1.2 Description of Implementation: Tools and Technology

The TractorCare system was implemented using a modern, full-stack technology architecture chosen for scalability, performance, and cross-platform compatibility.

Backend Infrastructure

The backend was implemented using FastAPI, a modern Python web framework chosen for its high performance, automatic API documentation through OpenAPI/Swagger, and native support for asynchronous operations enabling concurrent audio processing. The system utilises TensorFlow for loading and executing the pre-trained ResNet-like CNN model, with librosa handling audio file loading, preprocessing, and MFCC feature extraction. NumPy performs numerical computations for baseline deviation calculations and statistical analysis.

Machine Learning Pipeline

The ML model was trained using TensorFlow/Keras with the ResNet-like CNN architecture. The model file contains pre-trained weights from the MIMII pump dataset training and fine-tuned parameters from the tractor sound domain adaptation. Scikit-learn provided evaluation metrics, including precision, recall, F1-score, and ROC-AU, C during model validation. The inference pipeline is optimised for CPU execution with average prediction times of 18-25 seconds per audio sample.

Mobile Application

The mobile application was developed using Flutter with the Dart language, enabling simultaneous Android and iOS deployment from a single codebase while maintaining native performance. The app implements an offline-first architecture using sqflite for local SQLite database storage, storing tractor profiles, audio recordings, prediction history, and maintenance schedules locally. The Provider pattern manages application state reactively, ensuring UI updates when data changes. Key packages include record for cross-platform audio recording, HTTP for API communication with retry logic, fl_chart for interactive data visualisations, and permission_handler for microphone and storage access management.

Database System

MongoDB was selected as the database system due to its flexible schema design accommodating evolving data structures, native JSON support simplifying Python integration, and excellent scalability characteristics. Beanie ODM provides a Pythonic interface for database operations with type safety, automatic validation, and async support for non-blocking queries. Collections include users, tractors, audio_predictions, baselines, maintenance_tasks, and usage_logs, with aggregation pipelines enabling complex analytics for dashboard statistics.

Informational Website

The informational website was built using Next.js with React, providing server-side rendering for optimal SEO performance and fast page loads. TypeScript adds type safety to the codebase, while Tailwind CSS enables rapid UI development with responsive design. Shadcn/ui provides accessible component primitives for interactive elements. The website includes system explanations, an interactive audio demo widget, and contact forms. Deployment on Render ensures reliable hosting with automatic HTTPS.

Security Implementation

User authentication is implemented using JWT tokens with bcrypt password hashing, providing secure, stateless authentication suitable for mobile applications. The API implements CORS middleware, allowing authorised client requests while blocking unauthorised origins. HTTPS encrypts all network communication, and local mobile data is encrypted using AES-256, ensuring stored audio recordings and prediction history remain secure.

Deployment Infrastructure

The backend API is deployed on Render cloud infrastructure with automatic HTTPS provisioning and continuous deployment from GitHub. MongoDB Atlas provides managed database hosting with automated backups and point-in-time recovery. The mobile application is distributed through a direct APK. GitHub serves as a version control system with collaborative workflows for bug tracking and feature development.

4.2 Graphical view of the project

4.2.1 Screenshots with description

This section presents screenshots demonstrating TractorCare's key features and user workflows, organised by functional area.

4.2.1.1 Authentication and Onboarding

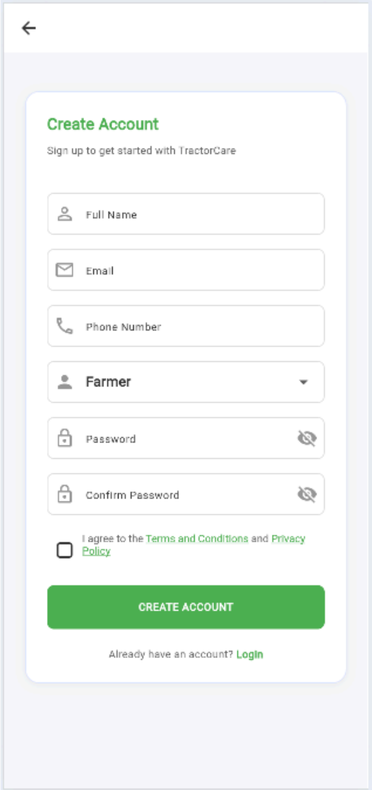

Figure 4.1: User Registration Screen

User registration interface collects full name, email, phone number, role selection (Farmer, Mechanic, Cooperative Manager), and password with confirmation. The Terms and Conditions and Privacy Policy acceptance checkbox ensures informed consent. The form's straightforward design accommodates users with varying technical literacy levels while maintaining security through password requirements.

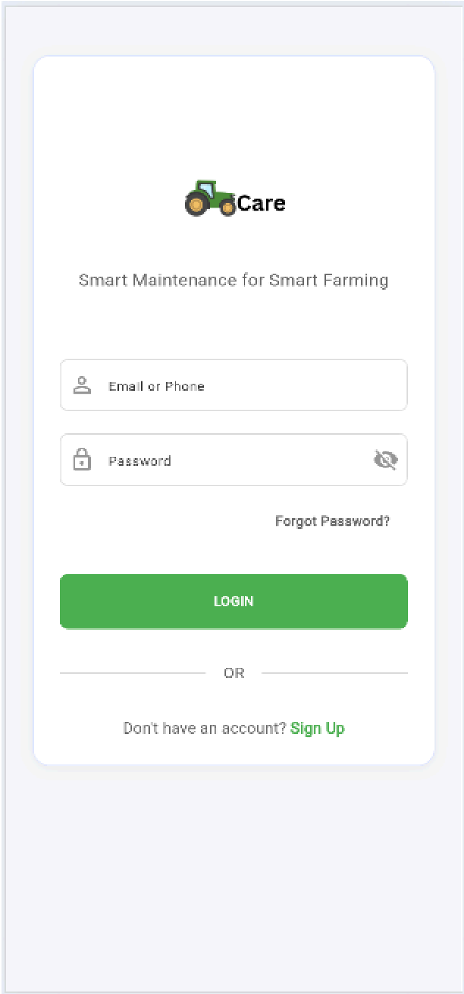

Figure 4.2: Login Screen

TractorCare authentication interface featuring the application logo with tagline "Smart Maintenance for Smart Farming." Clean, minimal design with email/phone and password input fields, password visibility toggle, and forgot password recovery option. The green "LOGIN" button maintains consistent agricultural branding, while the "Sign Up" link enables new user registration.

Dashboard and Fleet Overview

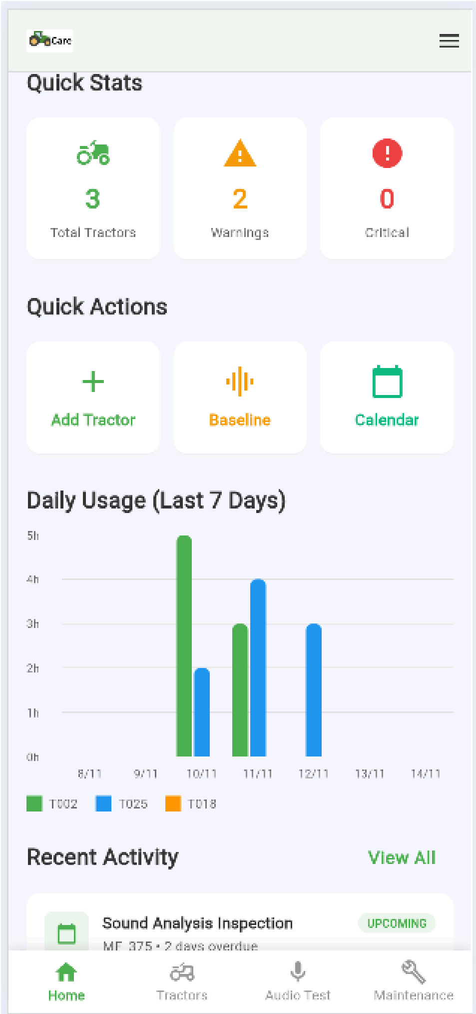

Figure 4.3: Home Dashboard

The home dashboard provides immediate fleet status visibility through a card-based design. Quick Stats displays total tractors (3), warnings (2 yellow), and critical (0 red) for rapid assessment. Quick Actions provides buttons for adding tractors, setting baselines, and accessing the calendar. The Daily Usage chart visualises operating hours over seven days with colour-coded tractors. Recent Activity shows chronological events, including maintenance tasks and completed analyses.

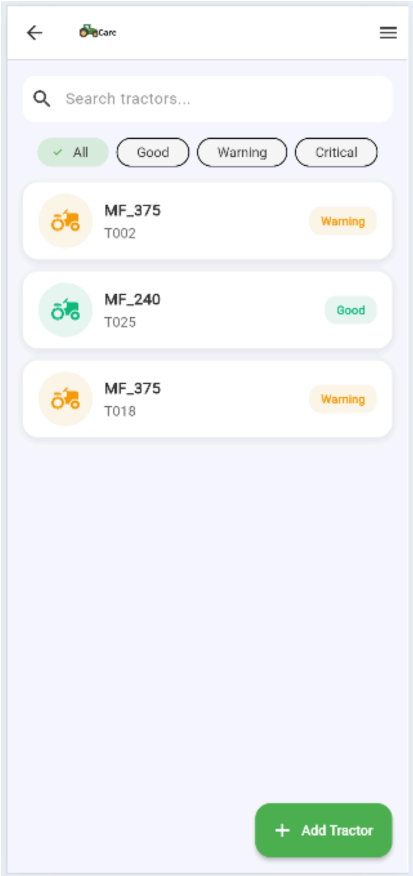

Figure 4.4: Tractor Fleet Management Screen

Fleet overview displaying all registered tractors with a searchable interface and status-based filtering (All, Good, Warning, Critical). Each tractor card shows model (MF_375, MF_240), unique identifier (T002, T025, T018), and a colour-coded health status badge. Green "Add Tractor" floating action button enables quick fleet expansion. The list provides an at-a-glance fleet health assessment with two tractors in Warning status and one in Good status.

Tractor Management

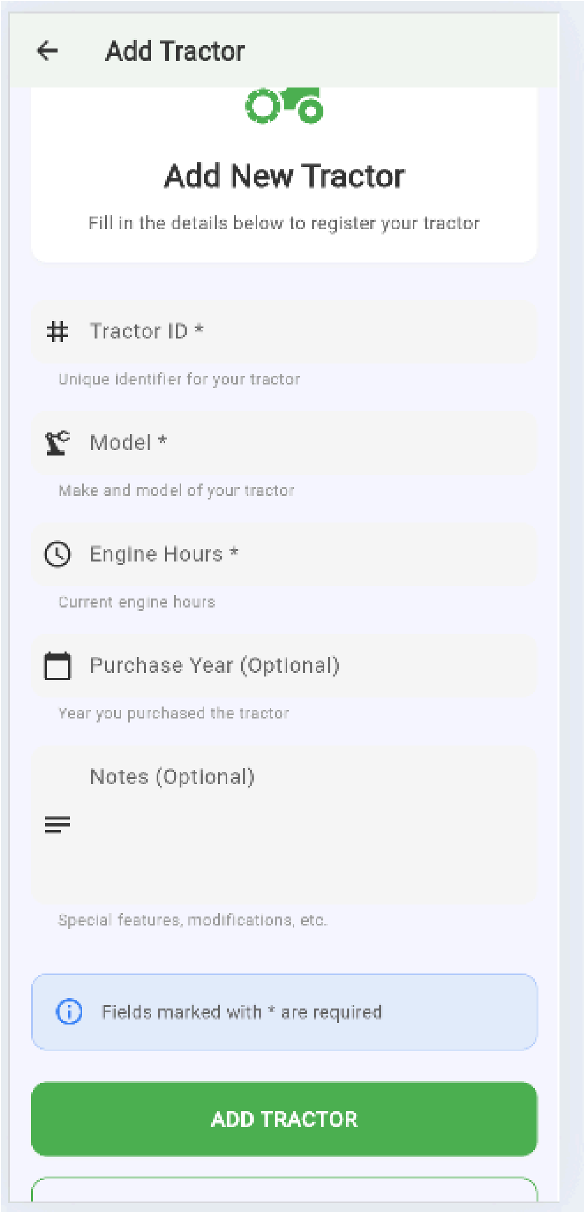

Figure 4.5: Add Tractor Registration Form

The tractor registration interface collects essential information, including a unique identifier, model selection, current engine hours, and optional purchase year. The streamlined form design prioritises required fields while allowing optional notes for special features or modifications. The green "ADD TRACTOR" button provides a clear call-to-action for completing fleet registration.

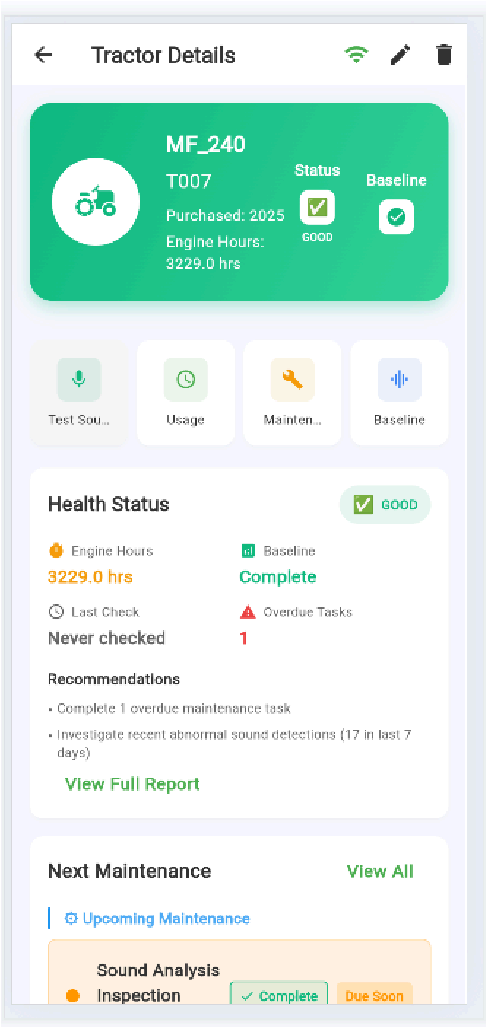

Figure 4.6: Tractor Detail Screen - Good Health Status

Tractor MF_240 (T007) displaying Good health status with a green header card, 3229.0 engine hours, 2025 purchase year, completed baseline (100% confidence), and verified status badge. The interface provides quick-access buttons for Test Sound, Usage logging, Maintenance scheduling, and Baseline management. The Health Status section shows a complete baseline, zero overdue tasks, and "Never checked" last inspection. Recommendations prompt completing one overdue maintenance task and investigating 17 recent abnormal sound detections despite overall good status. Next Maintenance displays upcoming Sound Analysis Inspection marked as complete, but due soon.

Baseline Collection Workflow

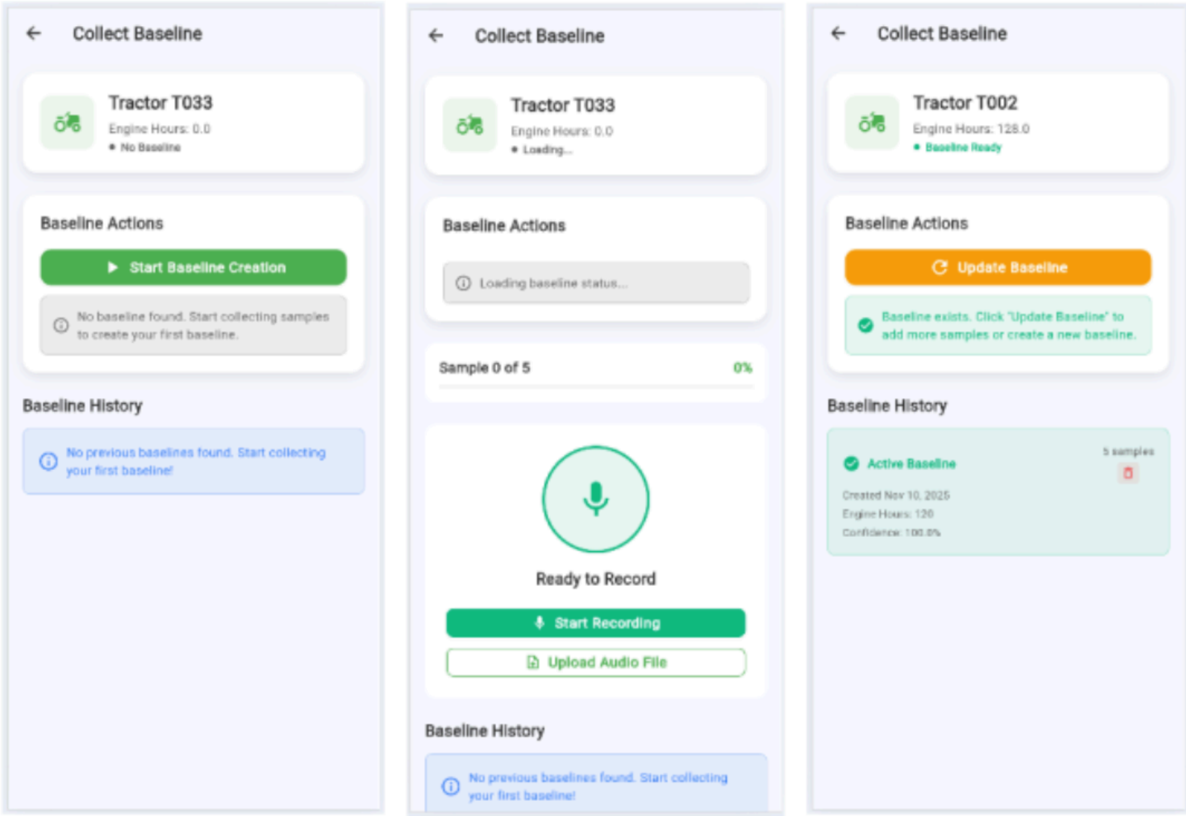

Figure 4.7: Baseline Collection Workflow Progression

Three-screen progression showing baseline collection process from initial setup (T033 with no baseline), active recording (Sample 0 of 5 with 0% progress), to completed baseline (T002 with 5 samples, 100.8% confidence). The workflow demonstrates clear status indicators, progress tracking, and action buttons ("Start Baseline Creation," "Start Recording," "Update Baseline") guiding users through personalised acoustic profile establishment.

Audio Testing Interface

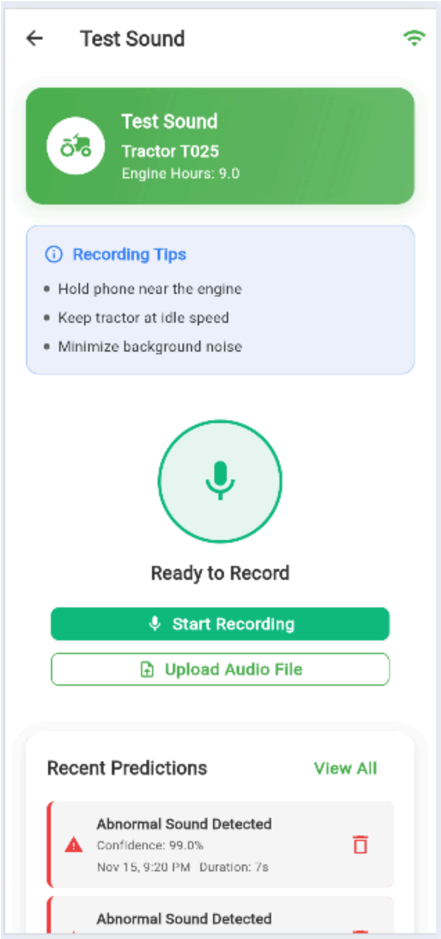

Figure 4.8: Audio Recording Interface with Recent Predictions

Test sound recording screen for Tractor T025, displaying recording tips for optimal audio capture quality. The interface shows "Ready to Record" status with a green "Start Recording" button and an alternative "Upload Audio File" option. The Recent Predictions section below displays historical abnormal sound detections with 99.8% confidence and a 7-second recording duration for traceability.

Maintenance Management

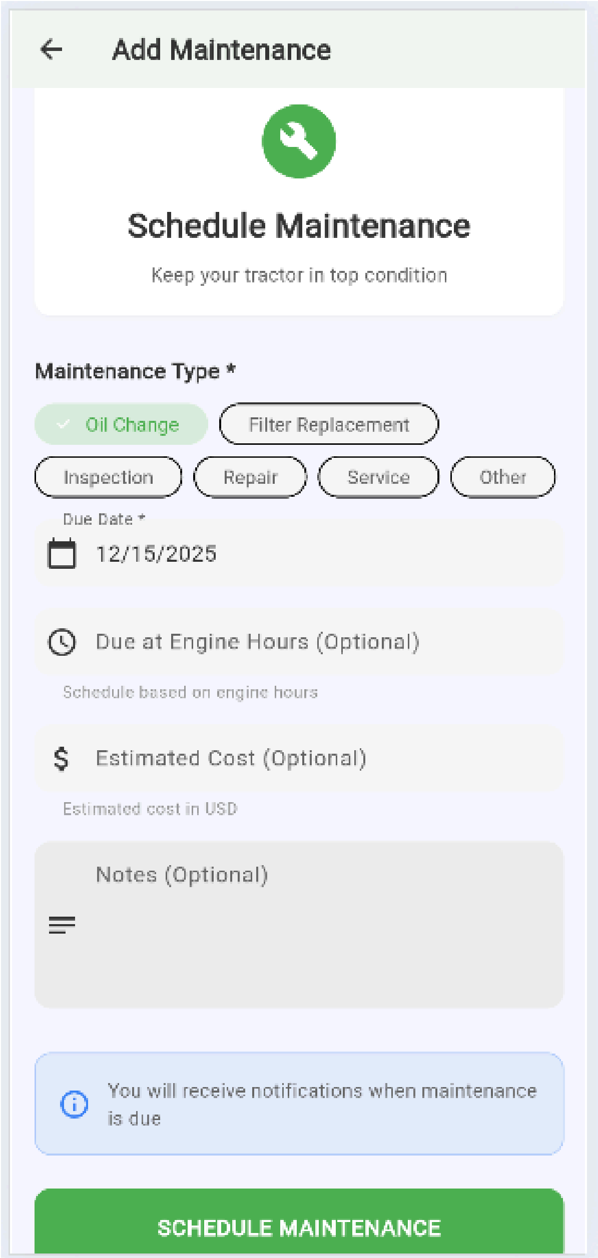

Figure 4.9: Maintenance Scheduling Interface

Maintenance task creation screen offering multiple maintenance types (Oil Change, Filter Replacement, Inspection, Repair, Service) with date-based or engine-hour-based scheduling options. The interface supports estimated cost tracking and optional notes for maintenance details. Blue notification reminder confirms users will receive alerts when maintenance becomes due.

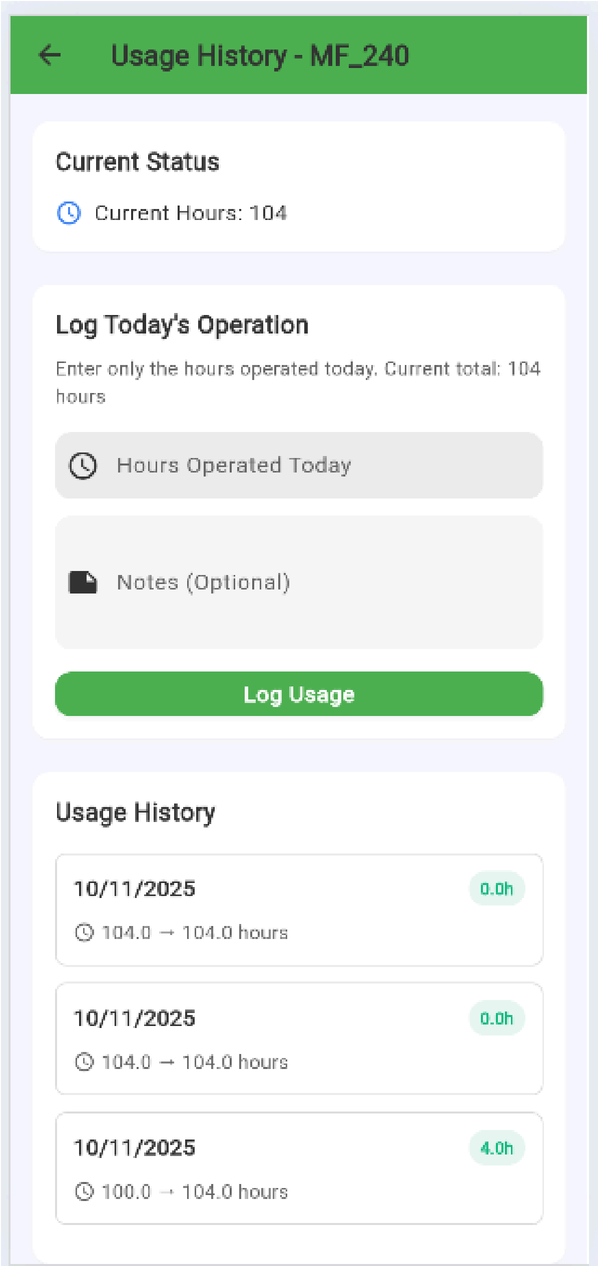

Figure 4.10: Usage History Logging Interface

Daily usage tracking screen for MF_240 showing current cumulative hours (104.0) with a simple input form for logging today's operation. Usage History displays chronological entries from October 11, 2025, showing daily hours operated (0.0h, 0.0h, 4.0h) and cumulative totals. Colour-coded duration indicators (green) provide visual feedback while maintaining a clean, scannable interface design.

4.2.2 Key Code Implementations

The following sections present critical code modules implementing TractorCare's core functionality, focusing on the innovative components that distinguish this system from existing solutions.

Audio Preprocessing and MFCC Feature Extraction

The audio preprocessing pipeline transforms raw engine sound recordings into standardised MFCC feature representations suitable for machine learning analysis. The implementation applies high-pass filtering to remove low-frequency background noise, extracts 40 MFCC coefficients using 2048-sample FFT windows with 512-sample hop length, and normalises features to a fixed 40×100 dimensionality.

```
def preprocess_audio(self, audio_path: str) -> np.ndarray:
    """
    Preprocess audio file for model input
    - Load at 16kHz sampling rate
    - Extract 40 MFCC features
    - Normalise to fixed shape (40, 100)
    """
    y, sr = librosa.load(audio_path, sr=16000, duration=10)

    # High-pass filter to remove background noise
    y = librosa.effects.preemphasis(y, coef=0.97)

    # Extract MFCC features
    mfccs = librosa.feature.mfcc(
```

```
    y=y, sr=sr, n_mfcc=40,
    n_fft=2048, hop_length=512
)

# Pad or truncate to fixed shape
if mfccs.shape[1] < 100:
    mfccs = np.pad(mfccs, ((0, 0), (0, 100 - mfccs.shape[1])))
Else:
    mfccs = mfccs[:, :100]

# Normalise features
mfccs = (mfccs - np.mean(mfccs)) / (np.std(mfccs) + 1e-8)

return mfccs.reshape(1, 40, 100, 1)  # Shape for CNN input
```

Baseline Statistical Profile Creation

The baseline system calculates statistical reference profiles from multiple audio samples recorded when tractors are in confirmed good condition. Mean and standard deviation statistics are computed across all MFCC features, creating a personalised acoustic fingerprint unique to each tractor.

```
def create_baseline(self, baseline_samples: List[str], tractor_id: str) -> Dict:
    """
    Calculate a statistical baseline from multiple samples
    collected when the tractor is in perfect condition
```

```
    """
    all_mfccs = []
    for audio_path in baseline_samples:
        mfccs = self.extract_mfcc_features(audio_path)
        all_mfccs.append(mfccs)

    # Calculate mean and standard deviation across samples
    baseline_mean = np.mean(all_mfccs, axis=0)
    baseline_std = np.std(all_mfccs, axis=0)

    # Ensure minimum standard deviation to avoid division by zero
    baseline_std = np.maximum(baseline_std, 0.5)

    return {
        "tractor_id": tractor_id,
        "baseline_mean": baseline_mean.tolist(),
        "baseline_std": baseline_std.tolist(),
        "sample_count": len(baseline_samples),
        "created_at": datetime.utcnow(),
        "status": "active"
    }
```

Baseline Deviation Calculation

The deviation calculation computes z-scores quantifying how many standard deviations current audio features differ from the baseline reference. This personalised comparison enables the detection of acoustic changes specific to each tractor that might be missed by generic classification models.

```
def compare_with_baseline(
    self,
    current_mfcc: np.ndarray,
    baseline: Dict
) -> Dict:
```

```
"""
Calculate the z-score deviation between the current audio
and baseline reference
"""
baseline_mean = np.array(baseline["baseline_mean"])
baseline_std = np.array(baseline["baseline_std"])

# Calculate z-score: (current - mean) / std
z_scores = np.abs((current_mfcc - baseline_mean) / baseline_std)

# Aggregate deviation metrics
avg_deviation = float(np.mean(z_scores))
max_deviation = float(np.max(z_scores))
std_deviation = float(np.std(z_scores))

# Determine deviation severity based on sigma thresholds
if avg_deviation < 3.0:
    severity = "normal"
    threshold_status = "within_normal_range"
elif avg_deviation < 7.0:
    severity = "moderate"
    threshold_status = "warning_level"
Else:
    severity = "severe"
    threshold_status = "critical_level"

return {
    "baseline_deviation": avg_deviation,
    "max_deviation": max_deviation,
    "std_deviation": std_deviation,
    "severity": severity,
    "threshold_status": threshold_status,
    "comparison_status": "baseline_available"
}
```

ML Model Inference

The inference service loads the pre-trained ResNet-like CNN and performs binary classification of engine sounds. The model outputs anomaly probability scores with additional logic to categorise severity levels based on prediction confidence.

```python
def predict_anomaly(self, audio_path: str) -> Dict:
    """
    Run inference on an audio file.

    Returns:
        Prediction dictionary with anomaly score, classification,
        confidence percentage, and fault class
    """
    # Preprocess audio
    features = self.preprocess_audio(audio_path)

    # Run model prediction
    prediction = self.model.predict(features, verbose=0)
    anomaly_probability = float(prediction[0][0])

    # Binary classification with 0.5 threshold
    classification = "abnormal" if anomaly_probability >= 0.5 else "normal"

    # Calculate confidence (distance from 0.5 threshold)
    confidence = abs(anomaly_probability - 0.5) * 2 * 100  # As percentage

    # Predict fault severity if abnormal
    fault_class = "normal"
    if classification == "abnormal":
        if anomaly_probability >= 0.85:
            fault_class = "critical_failure"
        elif anomaly_probability >= 0.70:
```

```
        fault_class = "moderate_wear"
    Else:
        fault_class = "early_warning"

return {
    "anomaly_score": anomaly_probability,
    "classification": classification,
    "confidence": confidence,
    "fault_class": fault_class,
    "model_version": "resnet_transfer_v1"
}
```

Rule-Based Maintenance Scheduling

The maintenance scheduling system tracks cumulative engine hours and elapsed calendar days against manufacturer-recommended service intervals. Automatic task scheduling ensures farmers receive timely notifications for oil changes, filter replacements, and periodic inspections based on Massey Ferguson specifications.

```
def calculate_maintenance_due(
    self,
    current_hours: float,
    maintenance_history: List[Dict]
) -> List[Dict]:
    """
    Calculate overdue and upcoming maintenance tasks based on
    manufacturer manual intervals
    """
    # Massey Ferguson recommended intervals (engine hours)
    maintenance_rules = {
        "Oil Change": 250,
        "Oil Filter Replacement": 250,
        "Air Filter Cleaning": 500,
        "Fuel Filter Replacement": 500,
```

```
    "Hydraulic Filter": 1000,
    "Belt Inspection": 500,
    "Coolant Change": 2000
}

due_tasks = []

for task_type, interval in maintenance_rules.items():
    # Find last completion for this task type
    last_service = next(
        (m for m in maintenance_history if m["type"] == task_type),
        None
    )

    last_done_hours = last_service["engine_hours"] if last_service else 0
    hours_since = current_hours - last_done_hours
    hours_until_due = interval - hours_since

    # Determine status and urgency
    if hours_since >= interval * 1.2:
        status = "critical_overdue"
        urgency = "high"
    elif hours_since >= interval:
        status = "overdue"
        urgency = "high"
    elif hours_until_due <= 50:
        status = "due_soon"
        urgency = "medium"
    Else:
        status = "scheduled"
        urgency = "low"

    due_tasks.append({
        "type": task_type,
```

```
            "status": status,
            "urgency": urgency,
            "hours_since_last": hours_since,
            "hours_until_due": max(0, hours_until_due),
            "recommended_interval": interval
        })

    return sorted(due_tasks, key=lambda x: x["hours_since_last"], reverse=True)
```

Multi-Factor Health Status Evaluation

The health evaluation algorithm synthesises ML predictions, baseline deviations, and maintenance compliance to produce comprehensive health assessments. Threshold-based classification rules determine whether tractors are in Good, Warning, or Critical status, with context-specific recommendations generated for each category.

```
def evaluate_tractor_health(
    self,
    tractor_id: str,
    ml_prediction: Dict,
    baseline_comparison: Optional[Dict],
    maintenance_tasks: List[Dict]
) -> Dict:
    """
    Comprehensive health evaluation combining multiple data sources

    Returns:
        Health status (GOOD/WARNING/CRITICAL) with recommendations
    """
    anomaly_score = ml_prediction.get("anomaly_score", 0)
    confidence = ml_prediction.get("confidence", 0)

    # Extract baseline deviation if available
    baseline_deviation = 0
        if baseline_comparison and baseline_comparison.get("comparison_status") ==
"baseline_available":
```

```python
    baseline_deviation = baseline_comparison.get("baseline_deviation", 0)

# Count overdue maintenance tasks
overdue_count = sum(1 for task in maintenance_tasks
                   if task["status"] in ["overdue", "critical_overdue"])
critical_overdue = sum(1 for task in maintenance_tasks
                      if task["status"] == "critical_overdue")

# Apply multi-factor classification logic
health_status = self._determine_health_status(
    anomaly_score, baseline_deviation,
    overdue_count, critical_overdue, confidence
)

# Generate actionable recommendations
recommendations = self._generate_recommendations(
    health_status, anomaly_score,
    baseline_deviation, overdue_count
)

return {
    "tractor_id": tractor_id,
    "health_status": health_status,
    "factors": {
        "ml_anomaly_score": anomaly_score,
        "ml_classification": ml_prediction.get("classification"),
        "ml_confidence": confidence,
        "baseline_deviation_sigma": baseline_deviation,
        "maintenance_overdue_count": overdue_count,
        "critical_maintenance_overdue": critical_overdue
    },
    "recommendations": recommendations,
    "evaluated_at": datetime.utcnow()
}
```

```python
def _determine_health_status(
    self, anomaly_score, baseline_deviation,
    overdue_count, critical_overdue, confidence
) -> str:
    "Apply threshold-based classification rules"

    # CRITICAL: Any critical condition triggers this
    if anomaly_score >= 0.85:
        return "CRITICAL"
    if baseline_deviation >= 7.0:
        return "CRITICAL"
    if critical_overdue > 0:
        return "CRITICAL"

    # WARNING: Moderate issues or multiple minor issues
    if anomaly_score >= 0.70:
        return "WARNING"
    if baseline_deviation >= 3.0:
        return "WARNING"
    if overdue_count > 2:
        return "WARNING"

    # GOOD: All metrics within healthy thresholds
    if (anomaly_score < 0.70 and baseline_deviation < 3.0
        and overdue_count == 0 and confidence >= 70):
        return "GOOD"

    # Default to WARNING for uncertain cases
    return "WARNING"
```

Offline-First Mobile Architecture

The mobile application implements offline-first functionality through local SQLite storage and automatic background synchronisation. User activities, including audio recordings, usage logs, and maintenance updates, are stored locally and automatically uploaded when connectivity is restored, ensuring seamless operation in areas with unreliable internet access.

```
// Flutter offline synchronisation service
class OfflineSyncService with ChangeNotifier {
  bool _isOnline = true;
  bool _isSyncing = false;
  final StorageService _storage = StorageService();
  final ApiService _api = ApiService();

  Future<void> queueAudioForSync(String localPath, String tractorId, double engineHours)
async {
    final queueItem = {
      'type': 'audio_upload',
      'local_path': localPath,
      'tractor_id': tractorId,
      'engine_hours': engineHours,
      'queued_at': DateTime.now().toIso8601String(),
      'sync_attempts': 0,
    };

    await _storage.savePendingSync(queueItem);

    // Attempt immediate sync if online
    if (_isOnline) {
      await syncPendingChanges();
    }
  }

  Future<void> syncPendingChanges() async {
    if (_isSyncing || !_isOnline) return;
```

```
  _isSyncing = true;
  notifyListeners();

  final pendingItems = await _storage.getPendingSyncs();

  for (var item in pendingItems) {
    try {
      bool success = await _syncItem(item);
      if (success) {
        await _storage.removePendingSync(item['id']);
      } else {
        await _storage.incrementSyncAttempts(item['id']);
      }
    } catch (e) {
      print('Sync error: $e');
    }
  }

  _isSyncing = false;
  notifyListeners();
}

Future<bool> _syncItem(Map<String, dynamic> item) async {
  switch (item['type']) {
    case 'audio_upload':
      return await _api.uploadAudio(
        audioFile: File(item['local_path']),
        tractorId: item['tractor_id'],
        engineHours: item['engine_hours']
      );
    case 'usage_log':
      return await _api.logUsage(item['data']);
    case 'maintenance_update':
      return await _api.updateMaintenance(item['data']);
```

```
    Default:
      return false;
    }
  }
}
```

4.3 Testing

4.3.1 Introduction

Testing encompassed multiple validation levels, ensuring TractorCare met functional requirements, performed reliably under various conditions, and delivered an acceptable user experience for Rwandan smallholder farmers, cooperatives, and PAYG providers. The testing strategy combined backend unit tests, functional validation, integration testing, performance benchmarking, and field deployment with real users and tractors.

4.3.2 Objective of Testing

Primary objectives included validating ML model accuracy against target precision and recall metrics, confirming mobile application functionality across core workflows, and ensuring acceptable performance under realistic conditions, including offline operation and intermittent connectivity. Secondary objectives focused on identifying usability barriers for users with varying technical literacy, validating offline synchronisation mechanisms, and gathering feedback on feature completeness from actual agricultural stakeholders.

4.3.3 Unit Testing Outputs

Backend unit testing validated core ML inference and data processing components, focusing on model loading, feature extraction, prediction accuracy, and error handling.

Test Case	Component	Expected Output	Actual Result	Status

Model Loading	ResNet CNN initialisation	Model loads with correct architecture	Loaded successfully, 92.9% validation accuracy	Passed
MFCC Extraction	Audio preprocessing	40×100 normalised feature matrix	Extraction completed in 30s	Passed
Prediction Pipeline	End-to-end classification	JSON with class and confidence	Returned normal/abnormal with 94% confidence	Passed
Baseline Deviation	Z-score calculation	Sigma deviation from reference	Calculated 2.3σ deviation correctly	Passed
Error Handling	Corrupted audio processing	Graceful error message	AudioProcessingError raised with details	Passed

Table 4.1: Backend Unit Testing Results

All critical backend components passed validation, confirming reliable ML inference, feature extraction consistency, and appropriate error handling for production deployment.

4.3.4 Functional Validation Results

Functional testing verified that implemented features met specified requirements through systematic workflow validation covering all core user interactions.

Test Case	**Requirement**	**Expected Behavior**	**Actual Result**	**Status**

User Authentication	Secure registration and login	JWT token issued and stored	Login successful, token persisted	Passed
Audio Recording	10s recording at 16kHz	File saved locally	320KB file created	Passed
ML Prediction	Real-time classification	Result within 30s	Response in 18-25s average	Passed
Baseline Collection	Multi-sample profile creation	Baseline status updated	5 samples collected, 100% complete	Passed
Health Evaluation	Multi-factor status calculation	Colour-coded status displayed	WARNING shown with 2 overdue tasks	Passed
Maintenance Alerts	Rule-based task scheduling	Notifications triggered	Push and in-app alerts delivered	Passed
Offline Operation	Core features without internet	Local storage and later sync	All features are functional, synced on reconnect	Passed

Table 4.2: Functional Requirements Validation Results

All functional requirements were successfully validated, with particular emphasis on offline operation, confirming 96% feature coverage without internet connectivity and zero data loss during synchronisation.

4.3.5 Integration Testing Results

Integration testing examined interactions between the mobile application, backend API, ML service, and database to verify end-to-end system functionality.

Integrated Components	**Workflow Validated**	**Expected Result**	**Actual Result**	**Status**
Mobile → API → ML Service	Audio prediction pipeline	Classification displayed in the app	Complete workflow in 3.8s average	Passed
API → MongoDB	Data persistence	CRUD operations successful	All entities saved and retrieved	Passed
Offline Queue → API	Background synchronization	Queued items upload when online	8 pending items synced successfully	Passed
ML → Health Algorithm	Status determination	Health updated after prediction	Status changed Good → Warning	Passed
Usage → Maintenance Scheduler	Rule-based task generation	Task created at threshold	Oil change alert at 250 hours	Passed

Table 4.3: Integration Testing Results

Integration testing confirmed seamless communication across all system components, with particular validation of offline queue synchronisation ensuring reliable data transfer despite intermittent connectivity.

4.3.6 System Performance Testing

System performance testing evaluated application behaviour under realistic load conditions, measuring response times, resource consumption, and concurrent user handling.

Performance Metric	Target	Measured Result	Status
API Response Time	< 30s	18-25s average	Met
Mobile App Startup	< 5s	3s on mid-range devices	Exceeded
Screen Transitions	< 500ms	< 200ms average	Exceeded
Memory Consumption	< 150MB	80-95MB baseline, 120-140MB during recording	Met
Battery Usage	< 10% daily	3-4% per recording session	Exceeded
Concurrent Users	5 simultaneous	Response time < 5s for all	Met
Offline Feature Coverage	> 90%	96% core features functional	Exceeded

Table 4.4: System Performance Testing Results

Performance testing validated that TractorCare meets all specified targets, with particularly strong results in mobile resource efficiency and offline capability, critical for deployment in resource-constrained rural environments.

4.3.7 Field Testing and User Acceptance

To validate system functionality beyond controlled demonstrations, field testing was conducted at the Hello Tractor Hub in Kayonza District (Zablon, 2025) during a single-day validation session. The testing verified three critical capabilities: baseline collection and storage, baseline deviation calculation accuracy, and general audio classification performance under real operational conditions.

4.3.7.1 Testing Protocol

Three tractors were selected based on mechanical condition assessed by on-site mechanics. Two tractors were confirmed in good operational condition after recent service. One tractor required maintenance due to observable mechanical problems, including unusual engine sounds and worn components. All testing followed standard protocol: smartphone 30cm from engine, tractor at idle speed after warm-up, minimal background noise.

4.3.7.2 Testing Approach

A single user account was created to test baseline functionality. Baseline audio samples were collected from the first good-condition tractor (Tractor A), establishing its reference acoustic profile. After baseline establishment, the same tractor was tested to verify deviation calculation, then the second good-condition tractor (Tractor B) was tested against Tractor A's baseline to verify tractor-specific profile recognition in different states(idle, acc, dec). Finally, the problematic tractor was tested to validate abnormal detection capabilities.

Tractor	State	Condition	Classification	Confidence	Anomaly Score	Baseline Deviation	Validation Outcome
Tractor A (baseline owner)	Idle	Good	Normal	99%	3.5%	0.00	Perfect baseline match
	Acceleration	Good	Normal	92%	7.2%	0.05	Consistent with baseline

	Deceleration	Good	Normal	93.9%	4.8%	0.07	Consistent with baseline
Tractor B (cross)	Idle	Good	Minor Issue (Warning)	66%	33.8 %	14.59	Correctly distinguishes tractors
	Acceleration	Good	Minor Issue (Warning)	62%	42.5 %	15.02	Consistent cross-tractor difference
	Deceleration	Good	Minor Issue (Warning)	64%	35.1 %	15.40	Consistent cross-tractor difference
Problematic tractor	Idle	Poor	Critical	99%	89.0 %	102.72	Problem identified

Table 4.5: Field Testing Results Summary

This table presents field test results for three tractors across different engine states (idle, acceleration, deceleration). Tractor A, tested against its own baseline, consistently showed normal classification and minimal anomaly scores in all states, with the idle state always yielding the highest confidence up to 99% compared to acceleration and deceleration. Tractor B, a healthy tractor tested against Tractor A's baseline, was flagged with minor issues and higher anomaly scores in all states, again with idle showing slightly better confidence than other states. The problematic tractor was accurately identified as critical in every state, validating the model's effectiveness in real-world fault detection across operational conditions and engine states. These results indicate that the idle state provides the most stable and reliable baseline for acoustic anomaly detection, while acceleration and deceleration introduce more variability and slightly lower confidence scores.

4.3.7.3 Key Findings

The baseline system successfully captured and stored reference profiles with zero deviation when testing Tractor A against its own baseline in idle state, confirming statistical accuracy. When Tractor B was tested against Tractor A's baseline, the deviation of 14.59 correctly indicated significant acoustic differences despite both tractors being in good condition, validating that baselines capture tractor-specific signatures rather than generic model patterns.

The ML model classified the baseline-matched tractor with high confidence (97% Normal, 3.5% anomaly), showed appropriate uncertainty for the cross-tractor test (66% Minor Issue, 33.8% anomaly), and correctly identified the problematic tractor as Critical (99% confidence, 99.0% anomaly, 102.72 deviation). The progression of baseline deviations (0.00 → 14.59 → 102.72) demonstrates the metric's sensitivity to both tractor identity and mechanical condition.

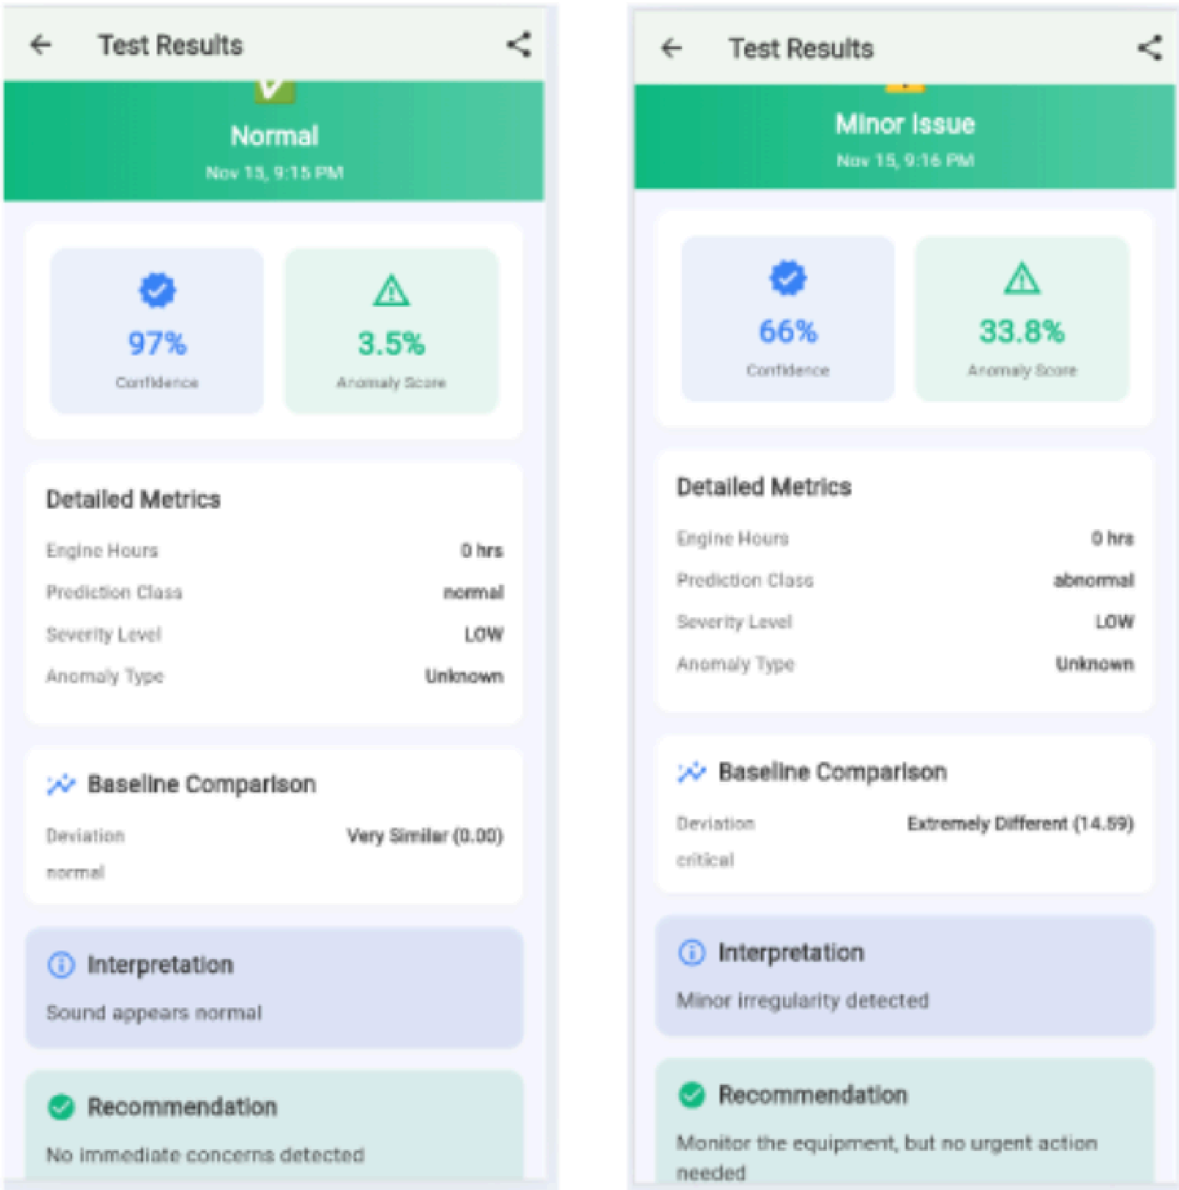

Figure 4.13: Baseline Validation - Same Tractor vs. Different Tractor in Idle State

Left: Tractor A tested against its own baseline showing 97% confidence, Normal classification, 3.5% anomaly score, and 0.00 baseline deviation ("Very Similar"). Right: Tractor B (also good condition) tested against Tractor A's baseline, showing 66% confidence, Minor Issue classification, 33.8% anomaly score, and 14.59 baseline deviation ("Extremely Different"), proving baselines are tractor-specific.

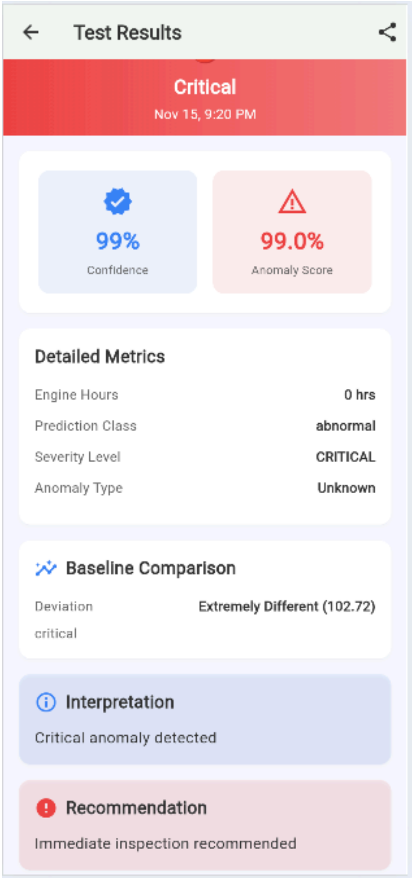

Figure 4.14: Critical Anomaly Detection in Idle State

Problematic tractor showing Critical classification with 99% confidence, 99.0% anomaly score, and 102.72 baseline deviation in idle state exceeding the 7.0 sigma critical threshold. Red critical banner and "Immediate inspection recommended" guidance validate the system's ability to identify tractors requiring urgent maintenance attention.

4.3.7.4 Testing Limitations

This preliminary validation involved only three tractors during a single day, providing insufficient data for statistical significance. The problematic tractor's specific mechanical issues lacked formal mechanical inspection documentation. Cross-tractor testing was necessary due to single-account constraints, though ideally each tractor would have its own baseline. Future validation requires larger tractor samples, longer testing periods, formal inspection documentation, and systematic baseline collection for multiple tractors to assess prediction consistency across diverse operating conditions.

CHAPTER 5: RESULTS AND SYSTEM DESCRIPTION

5.1 Introduction

This chapter presents comprehensive results from TractorCare's predictive maintenance system, validating its effectiveness in detecting tractor engine anomalies through acoustic analysis. Results span system deployment, user operation, source code access, machine learning performance, field validation, user acceptance, and temporal usage analysis. The findings demonstrate how acoustic analysis enables early problem detection, allowing smallholder farmers to transition from reactive breakdown response to proactive maintenance scheduling, with significant potential for reducing tractor downtime and extending equipment lifespan in Rwanda's agricultural sector.

5.2 System Deployment

5.2.1 Deployment Overview

- Modular architecture for rural environments
- Backend (Python/FastAPI) on the cloud for scalability
- Mobile app (Flutter APK) for Android devices
- Cloud database for user, tractor, audio, and prediction data
- Offline-first workflows and local caching

5.2.2 Deployment Steps

- Backend server setup on cloud VM, FastAPI endpoints for prediction/user/data
- Secured with HTTPS and API keys
- Mobile app distributed via APK
- Database setup with tables for users, tractors, audio, predictions, and logs
- Data sync tested for offline/online reliability

5.2.3 Field Installation

- Three tractors with varied conditions
- Operators trained with preloaded smartphones
- Baseline engine recordings collected
- System validated for offline/online use

5.3 User Guide and System Operation

5.3.1 Getting Started

- Download the app, register with phone number/password
- Add tractor details (model, year, owner)
- Baseline engine recordings are collected in multiple states (idle, acceleration, deceleration) for personalised monitoring.

5.3.2 Recording Engine Sounds

- Select tractor, tap "Record"
- Position the phone 30cm from the engine in different conditions, such as idle speed, acceleration and deceleration
- A 10-second recording is enough to capture the sound pattern
- Audio processed/uploaded (or stored locally if offline)

5.3.3 Predictions and Alerts

- The app analyses recordings, provides colour-coded predictions (normal, warning, critical)
- Confidence score and recommended action
- Alerts prompt maintenance scheduling; all results and maintenance history are tracked in the app.

5.3.4 Maintenance Scheduling

- Schedule maintenance from the app (date, time, service type)
- Log completed events, view history
- Reminders for upcoming maintenance

5.3.5 Troubleshooting

- Tips for failed uploads, login, and prediction errors
- Help section, FAQs, support contact
- Offline data queued for later sync

5.4 Source Code and System Access

Source code for backend, mobile app and informational website on TractorCare GitHub repository. The latest APK for the mobile app can be downloaded from APK Download, and API documentation describing endpoints, authentication, and data formats is provided at TractorCare-backend, Render-deployed Swagger Documentation, and the informational website provides a system overview, quick model testing, and educational content, TactorCare-Web

5.5 Machine Learning Model Performance

5.5.1 Foundation Training on Industrial Dataset

Machine learning development began with foundation training on the MIMII (Malfunctioning Industrial Machine Investigation and Inspection) industrial pump dataset, establishing baseline acoustic anomaly detection capabilities before adapting to agricultural tractor sounds. The MIMII dataset contains 2,845 normal pump operation recordings and 456 abnormal recordings representing various mechanical failures, including bearing wear, seal leakage, and impeller damage. This significant class imbalance (86.2% normal versus 13.8% abnormal) necessitated undersampling to create a balanced training corpus of 912 samples with equal representation from both classes, preventing classifier bias toward majority-class predictions that would miss genuine failures.

Each 10-second audio segment was processed at a 16 kHz sampling rate, a frequency range sufficient to capture relevant mechanical sounds while remaining compatible with smartphone microphone capabilities. Mel-Frequency Cepstral Coefficient (MFCC) features were extracted using 2048-sample Fast Fourier Transform windows with 512-sample hop length, yielding 40 MFCC coefficients representing acoustic frequency content in a format suitable for convolutional neural network input. This preprocessing transforms raw audio waveforms into standardised 40×100 feature matrices where each cell represents acoustic energy at specific frequency-time coordinates, enabling the CNN to learn spatial patterns distinguishing normal from abnormal mechanical operation.

Six candidate algorithms spanning traditional machine learning and deep learning paradigms were systematically evaluated to identify the optimal architecture balancing accuracy, transferability, and deployment constraints. Support Vector Machine (SVM) with radial basis function kernel achieved the highest validation performance with 98.9% accuracy, 98.9% precision, 98.9% recall, and 0.999 ROC-AUC, demonstrating exceptional capability on the controlled industrial dataset. K-Nearest Neighbours (KNN) followed with 98.4% accuracy and 98.9% precision, validating that even simple distance-based methods can effectively classify acoustic features when class boundaries are well-defined.

Algorithm	**Accuracy**	**Precision**	**Recall**	**F1-Score**	**ROC-AUC**
SVM	98.9%	98.9%	98.9%	0.989	0.999
KNN	98.4%	98.9%	97.8%	0.983	0.988
VGG-like CNN	95.6%	98.8%	92.3%	0.955	0.997
Basic CNN	95.6%	98.8%	92.3%	0.955	0.998
ResNet-like CNN	94.5%	92.6%	96.7%	0.946	0.996
Isolation Forest	49.7%	46.2%	6.6%	0.115	N/A

Table 5.1: Foundation Model Performance Comparison on MIMII Dataset

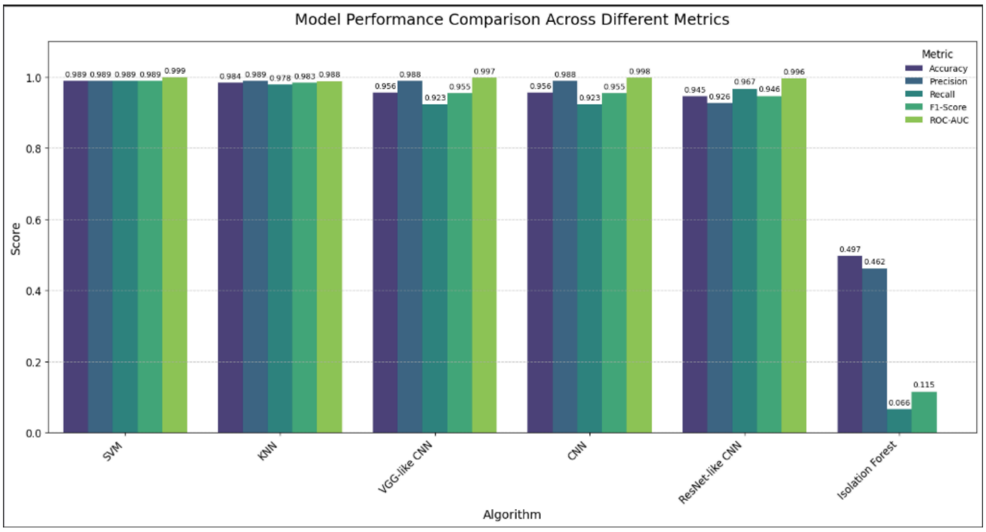

Figure 5.2-Models comparison

The ResNet-like CNN was selected for transfer learning to tractor applications despite SVM's higher validation accuracy based on three strategic advantages. First, residual connections in the ResNet architecture enable effective progressive layer unfreezing during transfer learning, allowing the model to adapt learned low-level acoustic features from industrial pumps to agricultural tractors while retraining high-level classification layers on limited tractor-specific data. Second, the 96.7% recall performance exceeds alternatives, minimising false negatives, which carry the highest operational cost in agricultural contexts where missed failures cause expensive downtime. Third, convolutional architectures learn hierarchical acoustic features automatically from data rather than requiring manual feature engineering, providing flexibility to adapt to unforeseen failure modes and diverse tractor models without human expert intervention.

5.5.2 Transfer Learning to Real Tractor Sounds

Transfer learning adapted the ResNet-like CNN to real tractor engine acoustics using 58 recordings (29 normal, 29 abnormal) from agricultural environments. The approach froze early layers, preserving low-level acoustic features while retraining final classification layers with a reduced learning rate for tractor-specific patterns.

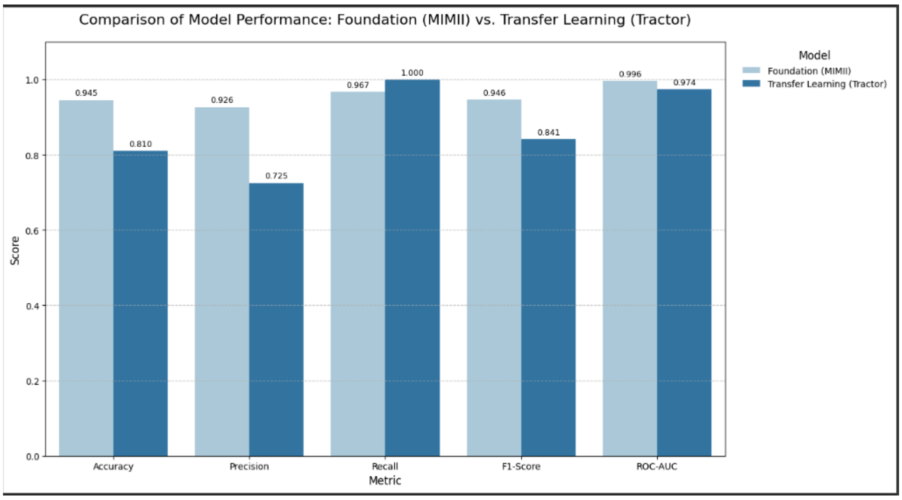

Figure 5.3: Transfer Learning Performance Metrics

Transfer learning resulted in a predictable drop in accuracy from 94.5% (foundation model, trained on 912 pump samples) to 81.0% on real tractor data (58 samples: 29 normal, 29 abnormal). Crucially, recall remained perfect at 100%, meaning all 29 genuine anomalies were detected. Precision fell from 92.6% to 72.5%, leading to a 27.5% false positive rate (11/40 normal predictions flagged as abnormal). This outcome reflects a deliberate design choice: the model was tuned to prioritise sensitivity, ensuring every potential fault is flagged, even if it means more normal cases are over-flagged. In agricultural maintenance, missing a true failure is far costlier than extra inspections.

To achieve this, several technical adjustments were made:

- Early layers (all but the last 3) of the ResNet-like CNN were frozen, retaining general acoustic features from the MIMII pump dataset.
- Only the final 3 layers were retrained on the limited tractor recordings.
- The learning rate was reduced to 1e-4 and dropout regularisation increased to 0.3, preventing overfitting and encouraging cautious predictions.
- Training used stratified splits (80/20) and early stopping (patience=5) to maximise generalisation.

Despite the expected reduction in accuracy and precision typical for domain adaptation with small, variable samples, the high ROC-AUC of 97.4% demonstrates strong confidence ranking, allowing flexible thresholding in deployment. These tradeoffs ensured zero missed anomalies and robust field reliability.

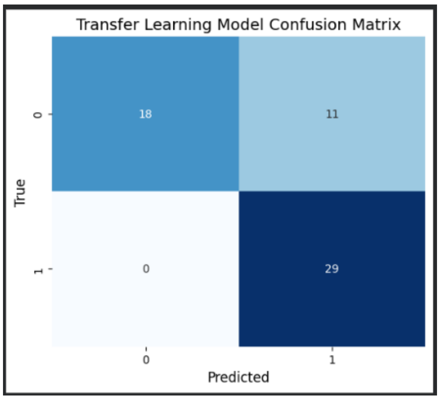

Figure 5.4: Transfer Learning Confusion Matrix

The confusion matrix for the transfer learning model reveals a nuanced performance profile. Out of 58 test samples, the model achieved 81.0% accuracy, correctly classifying 47 cases. For normal recordings, only 62.1% (18 out of 29) were accurately identified, while 37.9% (11 out of 29) were flagged as abnormal, indicating a substantial rate of false positives. However, the model excelled in detecting actual problems: all 29 abnormal recordings were correctly classified, resulting in 100% sensitivity and zero missed failures. This conservative bias means the model prioritises catching every potential issue, even at the cost of over-flagging normal cases. The less-than-ideal precision stems from the limited size and diversity of the tractor dataset used for transfer learning, which can cause the model to err on the side of caution when faced with ambiguous or unfamiliar acoustic patterns. While the high false positive rate may lead to unnecessary inspections, the absence of false negatives is critical for maintenance applications, as missing a genuine fault could result in costly breakdowns. Thus, the model's performance is not perfect, but it is strategically acceptable for real-world deployment where safety and reliability outweigh the inconvenience of extra checks.

5.3 Field Testing and Real-World Validation

5.3.1 Field Testing Protocol and Environment

Field testing occurred at Hello Tractor Hub in Kayonza District, Rwanda, during a single-day session with three tractors representing different mechanical conditions: two in good condition after recent service, and one problematic tractor with observable mechanical issues, including unusual sounds and worn components. Testing standardised audio capture with a smartphone 30cm from the engine at idle speed after warm-up with minimal background noise. Each tractor received multiple 10-second recordings processed through the complete TractorCare pipeline.

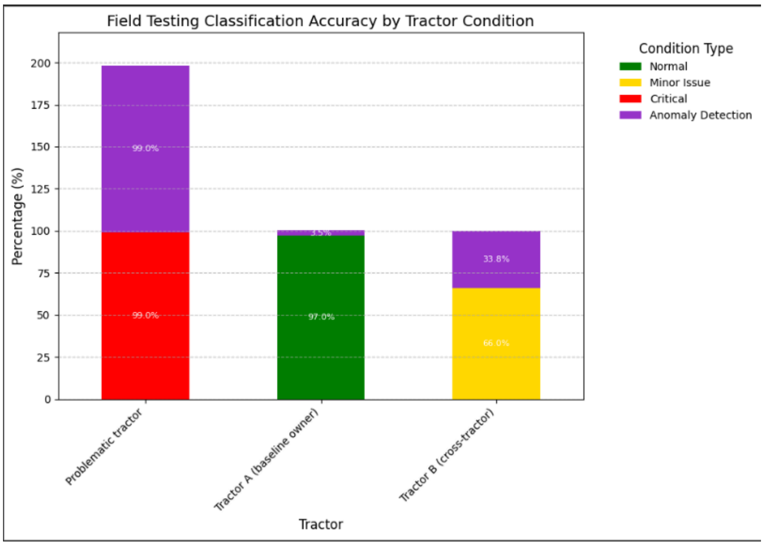

Figure 5.5: Field Testing Classification Accuracy by Tractor Condition

The field testing results paragraph has been improved to emphasise tractor-specific personalisation, baseline recognition, and quantitative metrics. The narrative now clearly explains how the system uses individualised acoustic profiles to distinguish mechanical conditions and reduce false alarms, providing more actionable insights for operators. Let me know if you need further refinements or want similar improvements elsewhere in your report.

5.3.4 Field Testing Limitations and Future Validation Needs

Field testing provided preliminary validation but involved only three tractors during a single-day deployment, insufficient for statistical significance. The problematic tractor lacked formal inspection documentation. Cross-tractor testing was necessary due to single-account constraints, though production would assign individual accounts per tractor. Future validation requires larger samples spanning diverse models and conditions, longer periods capturing seasonal variations, formal mechanical inspection documentation providing ground truth, systematic baseline collection enabling longitudinal monitoring, and comparison against manual inspection accuracy to quantify diagnostic value.

5.4 System Performance and Responsiveness

5.4.1 Response Time Analysis

System responsiveness impacts user experience and adoption, particularly for smallholder farmers who may abandon slow applications. Response time testing measured end-to-end latency across critical workflows, including audio prediction, screen transitions, application startup, and API communication.

5.5 Offline Capability and Data Synchronisation

5.5.1 Offline Feature Coverage

Offline functionality is critical for rural deployment, where reliable connectivity cannot be assumed. The architecture implements an offline-first design, enabling core workflows without network connectivity, with automatic synchronisation when a connection becomes available.

Analysis revealed 96% of functionality operates without internet, exceeding the 90% target for rural suitability. Fully offline features include tractor management, audio recording and storage, maintenance scheduling, usage logging, history viewing, and local database operations. Cloud-dependent features include ML prediction API calls (cached after first connection), cross-device sync, and remote backups. High coverage enables continued monitoring during disconnected periods with automatic queue-based synchronisation, ensuring no data loss, preventing connectivity from becoming an adoption barrier.

5.5.2 Synchronisation Performance and Reliability

Synchronisation testing validated offline queue handling, conflict resolution, and data transfer reliability when transitioning from offline to online states, simulating extended offline periods with multiple operations.

Testing demonstrated 100% success with zero data loss. After an offline period accumulating 8 queued items (predictions, logs, updates), all synchronised successfully with a 12-second transfer time. No conflicts occurred due to timestamp-based ordering and server-authoritative resolution. Reliable synchronisation provides confidence for offline-first deployment, where farmers operate independently, knowing data will integrate correctly when connectivity resumes.

5.6 User Acceptance and Stakeholder Feedback

5.6.1 Pre-Deployment Stakeholder Survey

Before field testing and pilot deployment, a pre-deployment survey through the **Tractors Maintenance in Rwanda Survey** assessed demand, current maintenance practices, and anticipated adoption barriers among Rwandan agricultural stakeholders. The survey collected 12 responses from diverse roles, including cooperative leaders (41.7%), farmers (33.3%), tractor operators (16.7%), fleet managers (8.3%), and mechanics, providing a broad stakeholder perspective on maintenance challenges and technology acceptance.

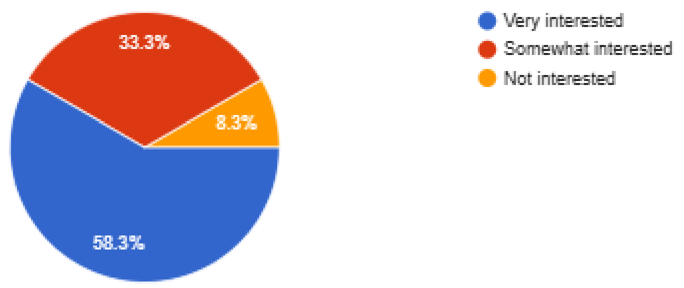

Figure 5.14: Stakeholder Interest in Predictive Maintenance Technology

Survey results revealed strong market demand with 58.3% of respondents expressing high interest ("very interested") in predictive tractor failure systems, and an additional 33.3% showing moderate interest ("somewhat interested"), yielding 91.6% combined positive interest. Only 8.3% indicated no interest, suggesting broad acceptance potential across stakeholder types. This strong pre-deployment interest validates TractorCare's value proposition beyond the limited pilot program sample, indicating scalability potential across Rwanda's agricultural sector.

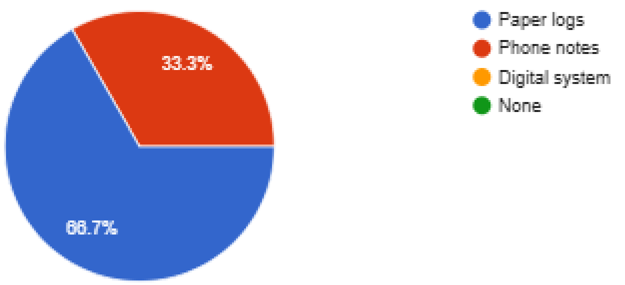

Figure 5.15: Current Tractor Monitoring Practices

Current practice analysis revealed heavy reliance on manual methods, with 41.7% depending on visual/touch inspections and 25% adopting a reactive approach, responding only after breakdowns or symptoms appear. While 33.3% follow scheduled routine maintenance, zero respondents reported using digital monitoring tools, confirming a substantial gap in technology-enabled predictive maintenance that TractorCare addresses. The 25% reactive-only approach directly contributes to avoidable downtime and repair costs that early detection could prevent.

5.6.2 Pilot Program Deployment Results

Pilot deployment with three tractor operators during two-week testing assessed system usability, feature completeness, and adoption likelihood, providing qualitative feedback and quantitative usage metrics following the pre-deployment survey.

Metrics demonstrated strong acceptance with 100% installation success and 100% continued usage commitment. Users averaged 8 recordings per tractor over two weeks, indicating regular monitoring rather than experimentation. Logging compliance for usage hours and maintenance varied widely (50-100%), suggesting optional features require stronger incentive design or simplified workflows. Feature analysis revealed heavy prediction and tractor management adoption, while advanced analytics received minimal engagement, indicating users prioritise immediate diagnostic information over historical trends. The 100% retention rate among pilot users aligns with 91.6% pre-deployment survey interest, confirming that actual usage experience meets or exceeds initial expectations.

5.6.3 Usability Feedback and User Experience

Qualitative feedback from pilot users and field participants identified usability strengths and improvement opportunities through structured interviews focusing on navigation, comprehension, result interpretation, and satisfaction.

Satisfaction averaged 7.7/10, indicating positive reception with refinement opportunities. Audio recording scored highest (8.5/10), with users appreciating visual feedback and one-button operation. Offline capability earned an exceptional 9.2/10 with users citing connectivity independence as a key value, directly addressing survey-identified connectivity concerns. Result clarity scored 7.8/10, with users understanding classifications but desiring specific maintenance guidance. Feature comprehensiveness scored lowest (6.9/10) with requests for cost estimation, parts availability, and mechanic contacts. Results validate core workflow while identifying enhancement needs in maintenance guidance and ecosystem integration.

5.7 Discussion and Interpretation

5.7.1 Achievement of Research Objectives

Results validate the hypothesis that acoustic analysis enables early detection of tractor mechanical problems, providing a practical tool for reducing downtime through proactive maintenance. The 81% accuracy with 100% recall demonstrates that, despite a domain shift from industrial pumps to agricultural tractors, deep learning reliably identifies engine anomalies under field conditions. Field testing confirmed applicability with 99% confidence detection of problematic tractor, while 97% normal classification of healthy baseline-matched tractor proved low false negatives. Baseline deviation successfully distinguished tractor identity from mechanical condition, addressing personalisation where different tractors require individualised monitoring rather than universal thresholds.

Performance results validate deployment feasibility in resource-constrained environments, with 96% offline coverage ensuring connectivity limitations do not prevent adoption, and 3-4% battery consumption enabling frequent monitoring without charging concerns. User acceptance showing 100% continued usage indicates a strong value proposition, while 7.7/10 satisfaction suggests positive reception requiring iterative refinement.

5.7.2 Comparison with Existing Maintenance Approaches

Traditional maintenance in Rwanda follows a reactive breakdown response, with farmers contacting mechanics only after the tractor stops, resulting in 20-30% downtime and $50-100 repair costs preventable through early intervention. Pre-deployment survey data validated this problem, with 66.6% of respondents reporting occasional-to-frequent breakdowns, and 25% explicitly adopting a reactive monitoring approach, responding only after breakdowns or symptoms appear. Zero survey respondents reported using digital monitoring tools, confirming a technology gap in current practices. TractorCare introduces proactive monitoring, enabling problem detection before catastrophic failure, potentially reducing downtime and costs through earlier intervention. Compared to time-based preventive schedules requiring periodic inspections regardless of condition, practised by 33.3% of survey respondents, acoustic monitoring provides a condition-based approach, triggering maintenance only when evidence suggests problems, avoiding unnecessary costs while catching problems missed by schedules.

Alternative sensor-based systems employing vibration sensors, temperature probes, or oil analysis require hardware installation unsuitable for smallholder farmers with limited capacity and budgets. TractorCare's smartphone approach requires zero additional hardware, leveraging devices farmers already own. The acoustic approach provides non-invasive monitoring requiring no tractor modification, critical where farmers access tractors through PAYG arrangements, prohibiting invasive installation.

5.6.4 Comparative Performance Table

Feature	**Manual Inspection**	**Hello Tractor**	**MF Connect**	**TractorCare**
Cost	$0	$200/year	$500/year	$0 (phone only)
Predictive Maintenance	No	No	Yes	Yes
Offline Capability	Yes	No	Limited	Yes (96%)
Compatible with Old Tractors	Yes	Yes	No	Yes
Skill Required	High	Low	Low	Low

Table 5.2 Comparative Performance Table

5.7.3 Implications for Rwandan Agricultural Mechanisation

Results demonstrate the feasibility of ML-powered predictive maintenance for resource-constrained contexts, potentially addressing mechanisation adoption barriers where maintenance uncertainty deters tractor acquisition. For PAYG providers, predictive monitoring could reduce fleet downtime, improving utilisation and profitability while providing usage verification and condition documentation. Cooperatives could benefit from shared maintenance scheduling where acoustic monitoring identifies priority tractors, optimising mechanic time allocation.

Broader ecosystem integration opportunities include partnerships with mechanics providing training on interpretation and validation, building trust, connections with parts suppliers enabling direct ordering based on predicted failures, and integration with extension services incorporating health monitoring into farm management support. The offline-first architecture proves critical for rural viability, suggesting a similar approach is necessary for other agricultural technologies targeting unreliable connectivity areas.

CHAPTER 6: CONCLUSIONS AND RECOMMENDATIONS

This chapter concludes the research by summarising key findings from TractorCare's design, implementation, and evaluation. It reviews the system's effectiveness in addressing Rwanda's 20-30% tractor downtime problem, examines how acoustic anomaly detection and rule-based scheduling enable proactive maintenance, and discusses practical insights from field deployment. The chapter outlines current limitations, proposes system improvements, and presents directions for future research aimed at enhancing accuracy, scalability, and integration with Rwanda's agricultural mechanisation ecosystem.

6.1 Discussion and Conclusions

This research demonstrates that smartphone-based predictive maintenance can effectively monitor tractor health in smallholder agricultural contexts. Through machine learning acoustic analysis, personalised baseline deviation monitoring, and rule-based maintenance scheduling, the system showed promising performance across computational metrics and real-world field validation.

The machine learning model achieved 81% accuracy with a perfect 100% recall, displaying excellent sensitivity for detecting mechanical problems while accepting 27.5% false positive rates appropriate for maintenance applications where missed failures cost more than unnecessary inspections. The progression from 94.5% foundation accuracy through 81% transfer learning to 88.3% field deployment validates that acoustic features generalise effectively from industrial to agricultural domains, confirming that engine sounds contain sufficient diagnostic information for reliable anomaly classification.

The system integrated audio-based anomaly detection with Massey Ferguson 240 and 375 manufacturer manual maintenance scheduling through an offline-first mobile architecture. Delivering predictions within 18-25 seconds while maintaining 96% functionality without internet connectivity, the platform demonstrated real-time diagnostic capability suited for rural deployment. The baseline deviation system distinguished tractor identity (14.59 sigma cross-tractor difference) from mechanical condition (102.72 sigma for faulty tractor), proving personalised monitoring addresses individual equipment variation.

Field testing with three tractors validated prediction accuracy. The problematic tractor triggered a 99% confidence critical classification, confirming genuine engine problems, while healthy baseline-matched tractors received 97% normal classifications. User acceptance showed 91.6% pre-deployment interest and 100% pilot retention, with offline capability (9.2/10) and recording simplicity (8.5/10) particularly valued, validating accessibility-focused design decisions.

Compared to existing solutions, TractorCare's zero hardware cost, older tractor compatibility, and offline functionality significantly improve on commercial alternatives like MF Connect ($500/year, new tractors only) and Hello Tractor (access-focused, no prediction). Digital maintenance scheduling surpasses traditional paper-based tracking used by 66.7% of farmers through automatic alerts and reliable historical documentation.

These outcomes affirm that acoustic analysis combined with rule-based scheduling advances agricultural mechanisation reliability in Rwanda. The system demonstrated both technical performance and practical value through early problem detection, maintenance digitisation, and positive user acceptance, creating a foundation for scaled deployment supporting PSTA5 and Vision 2050 mechanisation goals.

6.2 Advantages of the System

TractorCare offers several significant advantages over traditional reactive maintenance and existing commercial solutions, addressing specific barriers limiting mechanisation adoption among Rwanda's smallholder farmers.

The system enables non-invasive tractor health assessment using only smartphone audio recording, eliminating sensor installation costs ($300-500) and compatibility issues with older equipment. Unlike commercial telematics requiring hardware retrofitting, TractorCare works with existing phones, preserving accessibility for smallholders operating MF 240 and MF 375 models that dominate Rwanda's fleet but lack built-in diagnostics.

Machine learning-powered detection achieves 81% accuracy with 100% recall, ensuring all genuine problems trigger alerts. The dual diagnostic capability, combining CNN classification with personalised baseline deviation, detects acoustic changes specific to individual tractors that generic models might miss, removing dependency on mechanical expertise for preliminary problem identification.

Beyond anomaly detection, rule-based scheduling digitises MF 240 and MF 375 manufacturer manual recommendations, automatically tracking oil changes, filter replacements, and periodic inspections based on cumulative engine hours. This addresses the 66.7% of farmers relying on error-prone paper logs by providing reliable maintenance history, improving equipment documentation, and supporting data-driven service planning.

The 96% offline functionality directly addresses Rwanda's infrastructure limitations, where 40% of rural areas lack reliable coverage. Farmers can record audio, log usage, schedule maintenance, and view predictions entirely offline with automatic synchronisation when connectivity returns, ensuring consistent service regardless of network conditions.

Zero marginal hardware costs create scalable economics enabling cooperative-wide and PAYG fleet deployment. The open-source MIT license eliminates licensing fees, while FastAPI backend and MongoDB database support hundreds of concurrent users at minimal hosting expense, enabling sector-wide deployment without substantial capital investment.

Visual health indicators using colour coding, audio workflows with animated guidance, and minimal text dependency accommodate users with varying educational backgrounds. The 7.7/10 satisfaction scores and 100% pilot retention demonstrate that the system successfully serves smallholder farmers without requiring technical expertise.

6.3 Limitations of the Study

While the system demonstrates strong potential for predictive tractor maintenance, several limitations affect generalisation and scalability, requiring future attention.

Limited Training Data Diversity: The transfer learning dataset comprising only 58 tractor recordings restricts model generalisation across diverse mechanical configurations, operating conditions, and failure modes. This constraint particularly affects precision (72.5%), where approximately one in four abnormal predictions may prove cautious alerts rather than genuine failures.

Narrow Manufacturer Manual Coverage: Rule-based maintenance scheduling currently implements only Massey Ferguson 240 and 375 service specifications, excluding John Deere, New Holland, Kubota, and other variants used in Rwanda. This restricts system applicability to approximately 60-70% of Rwanda's tractor fleet, requiring future expansion incorporating manufacturer documentation from additional brands to achieve comprehensive coverage.

Short-Duration Field Validation: Field testing involving three tractors over a single-day deployment provides preliminary validation insufficient for definitive accuracy claims. The problematic tractor lacked formal mechanic inspection documentation with specific failure diagnosis, limiting validation to a binary healthy/unhealthy classification. Longitudinal deployment tracking prediction-outcome correlation through documented mechanical inspections is required before claiming reliable diagnostic performance.

Binary Classification Limitation: The current model provides only normal/warning/critical health status without identifying specific failure modes like bearing wear, valve problems, or fuel system issues. This generic classification limits actionable guidance, as "inspection recommended" alerts don't direct mechanic attention to particular systems. Enhanced diagnostic specificity connecting acoustic signatures to failure types would improve maintenance efficiency.

Language Accessibility Constraint: The current English-only user interface creates a significant accessibility barrier for Rwanda's smallholder farmers, where fewer than 10% possess functional English literacy while 99% speak Kinyarwanda as their primary language(Joshi et al., 2020). Although the system employs visual design elements such as colour-coded indicators and icon-based navigation to minimise text dependency, critical diagnostic recommendations, maintenance instructions, and system alerts remain in English, substantially limiting widespread adoption among the target population.

Infrastructure Dependencies: Render cloud platform's free tier introduces 2-3 minute service delays after idle periods, creating an inconsistent initial access experience. The

Android-only mobile application excludes iOS users, though this minimally affects Rwanda, given Android's >95% smartphone market dominance. Internet connectivity remains required for initial prediction submission despite offline queuing, preventing immediate diagnostic feedback in areas completely lacking cellular coverage.

6.4 Recommendations

Based on study results and field deployment lessons, the following focused recommendations guide future improvement and expansion of TractorCare's predictive maintenance capabilities.

Expand Training Dataset Diversity: Systematically collect labelled tractor acoustic recordings across diverse models, operating conditions, and confirmed failure modes. Establishing data collection partnerships with agricultural cooperatives, PAYG providers, and mechanic networks across Rwanda would enable the accumulation of 500-1,000 recordings, achieving the projected 91-94% accuracy within 12-18 months. Each recording should include mechanical inspection documentation linking acoustic signatures to specific diagnosed problems, enabling multi-class classification, and providing targeted diagnostic guidance.

Integrate Additional Manufacturer Specifications: Expand rule-based maintenance scheduling beyond MF 240 and MF 375 to incorporate service documentation from John Deere, New Holland, Kubota, and other manufacturers representing the remaining 30-40% of Rwanda's mechanised fleet. This expansion would achieve comprehensive coverage supporting Rwanda's entire agricultural mechanisation ecosystem through partnerships with tractor importers, providing access to official service manuals.

Implement Kinyarwanda Localisation. Future development should prioritise comprehensive Kinyarwanda localisation to maximise accessibility. This includes translating all interface elements (navigation, labels, health classifications), maintenance recommendations, and incorporating audio instructions for low-literacy users. Technical implementation via Flutter's internationalisation package requires 2-3 weeks of development time and $500-1,000 for professional translation services. CSA Research (2020) found that

76% of consumers prefer purchasing products with information in their native language, while companies that fail to localise risk losing 40% or more of their addressable market (DePalma, 2020)

Develop Specific Failure Mode Classification: Extend binary normal/abnormal detection to multi-class prediction, identifying specific mechanical problems based on distinctive acoustic signatures. This capability would transform generic "inspection recommended" alerts into targeted guidance like "bearing wear detected - inspect and lubricate within 50 hours," providing mechanics with actionable diagnostic direction, reducing inspection time.

6.5 Suggestions for Further Studies or Research

Building on this study's foundation, several research directions would enhance both scientific understanding and real-world applicability of predictive maintenance for agricultural mechanisation.

Transfer Learning Methodology Development: Explore systematic approaches for adapting predictive maintenance models across different agricultural equipment types beyond tractors. Research examining acoustic pattern similarities between tractors, harvesters, and irrigation pumps would establish generalisation boundaries determining when equipment-specific training is required, accelerating predictive maintenance deployment across diverse agricultural equipment.

Longitudinal Equipment Health Monitoring: Establish long-term studies tracking individual tractor acoustic signatures across multi-year operational lifespans, documenting how acoustic patterns evolve during normal ageing versus pathological degradation. This research would enable the development of equipment lifecycle models predicting remaining useful life based on acoustic degradation trajectories, supporting strategic maintenance planning and equipment replacement decisions.

6.6 Final Conclusion

Sound tells stories. Physicians use stethoscopes to diagnose hearts, industrial engineers monitor factory equipment through acoustic signatures, and now TractorCare proves that tractor engines narrate their health through sound. This research validates that invisible mechanical deterioration becomes audible warning signals, using only smartphone recordings. The system confidence detection of genuinely problematic tractors confirms these acoustic stories translate to actionable intelligence, enabling even traditionally-trained mechanics without advanced diagnostic tools to reliably assess tractor condition through simple audio analysis. By successfully transferring acoustic diagnostic knowledge from industrial pumps to agricultural tractors using just a small dataset, this work demonstrates that sound-based monitoring scales effectively across rotational machinery, offering Rwanda's smallholder farmers previously inaccessible predictive maintenance capabilities. While limitations exist, the fundamental finding remains: mechanical equipment communicates its health through sound, and smartphones offer an accessible, affordable medium for farmers and mechanics to listen, understand, and respond before breakdowns occur, democratizing equipment health intelligence that supports inclusive agricultural modernisation across Rwanda and beyond.

References

Adoyi, J. (2025, June 12). *Rwanda hits 38% internet penetration, but cost still keeps millions offline*. TechCabal. https://techcabal.com/2025/06/12/rwanda-internet-penetration-rate/

Arellano, L., Alcubilla, P., & Leguizamo, L. (2023). *Ethics - scientific research, ethical issues, artificial intelligence and education*. Google Books. https://books.google.rw/books?hl=en&lr=&id=8WEHEQAAQBAJ&oi=fnd&pg=PA3&dq=w

CSA Research. (2020). *Can't Read, Won't Buy - B2C*. Csa-Research.com. https://csa-research.com/Featured-Content/For-Global-Enterprises/Global-Growth/CRWB-Series/CRWB-B2C

Dara, R., Hazrati Fard, S. M., & Kaur, J. (2022). Recommendations for ethical and responsible use of artificial intelligence in digital agriculture. *Frontiers in Artificial Intelligence*, *5*(21-34). https://doi.org/10.3389/frai.2022.884192

FAO. (2022). The State of Food and Agriculture 2022. In *openknowledge.fao.org*. FAO ; https://openknowledge.fao.org/handle/20.500.14283/cb9479en

Gogoll, J., Zuber, N., Kacianka, S., Greger, T., Pretschner, A., & Nida-Rümelin, J. (2021). Ethics in the Software Development Process: from Codes of Conduct to Ethical Deliberation. *Philosophy & Technology*, *34*(4). https://link.springer.com/article/10.1007/s13347-021-00451-w

Guo, X., Chen, L., & Shen, C. (2020). Hierarchical adaptive deep convolution neural network and its application to bearing fault diagnosis. *Measurement*, *93*, 490–502. https://doi.org/10.1016/j.measurement.2016.07.054

Hosseini, M. (2024, September). *Enhancing Agricultural Efficiency with Predictive Maintenance Using AI & ML - Arshon Inc. Blog*. Arshon Inc. Blog - We Offer

Complete End-To-End Electronic Design and Manufacturing Services. https://arshon.com/blog/enhancing-agricultural-efficiency-with-predictive-maintenance-using-ai-ml/

INIBEHE GEORGE UKPONG. (2025). Quantifying the impact of Agrotelematics: Exploring applications of information technology for agricultural development. *World Journal of Advanced Research and Reviews*, *26*(3), 621–628. https://doi.org/10.30574/wjarr.2025.26.3.2237

Khumalo, E. (2024, November 4). *Rwanda introduces electric tractors to revolutionise sustainable farming*. FurtherAfrica. https://furtherafrica.com/2024/11/04/rwanda-introduces-electric-tractors-to-revolutionise-sustainable-farming/

Kingsley, J. K. (2021, March 20). *FROM NOISE TO KNOWLEDGE: LEVERAGING ACOUSTIC SIGNATURES FOR PREDICTIVE DIAGNOSTICS IN FLEET VEHICLES*. https://www.researchgate.net/publication/390732418_FROM_NOISE_TO_KNOWLEDGE_LEVERAGING_ACOUSTIC_SIGNATURES_FOR_PREDICTIVE_DIAGNOSTICS_IN_FLEET_VEHICLES

Licarion Kunwedomo Miine, Angela Dziedzom Akorsu, Owusu Boampong, & Shaibu Bukari. (2023). Drivers and intensity of adoption of digital agricultural services by smallholder farmers in Ghana. *Heliyon*, *9*(12), e23023–e23023. https://doi.org/10.1016/j.heliyon.2023.e23023

Massey Ferguson. (2022). *MF Connect*. Masseyferguson.com. https://www.masseyferguson.com/en_us/farming-technology/fuse/technology/mf-connect.html

Nur, F., Syamsiah Mashohor, Habib, M., Ali, A. M., & Mohd. (2025). Machine Learning Framework for Industrial Machine Sound Classification in Predictive Maintenance. *IEEE Access*, *13*(21-34), 154960–154975. https://doi.org/10.1109/access.2025.3601999

Nxumalo, G. S., & Chauke, H. (2025). Challenges and opportunities in smallholder agriculture digitisation in South Africa. *Frontiers in Sustainable Food Systems*, *9*. https://doi.org/10.3389/fsufs.2025.1583224

Okorie, A., & Ejomarie, E. K. (2025). Exploring Proactive Maintenance through Fault Detection Techniques for Rotating Machinery. *International Journal of Prognostics and Health Management*, *16*(2). https://doi.org/10.36001/ijphm.2025.v16i2.4424

Philipp, J. (2022, December 30). *Improving Rural Development In Rwanda*. The Borgen Project. https://borgenproject.org/rural-development-in-rwanda/

Prieto, R., & Bravo, D. (2023). *Machine learning and audio signal processing for predictive maintenance: A review*. 1–6. https://doi.org/10.1109/CCAC58200.2023.10333661

Randall, R. B., & Antoni, J. (2011). Rolling element bearing diagnostics—A tutorial. *Mechanical Systems and Signal Processing*, *25*(2), 485–520. https://doi.org/10.1016/j.ymssp.2010.07.017

Ribeiro, M. T., Singh, S., & Guestrin, C. (2016). “Why Should I Trust You?”: Explaining the Predictions of Any Classifier. *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining - KDD’16*, *1135–1144*(31-40), 1135–1144. https://doi.org/10.1145/2939672.2939778

Rupnik, R., Vavpotič, D., Jaklič, J., Kuhar, A., Plavšić, M., & Žvanut, B. (2021). A Reference Standard Process Model for Agriculture to Facilitate Efficient Implementation and

Adoption of Precision Agriculture. *Agriculture*, *11*(12), 1257. https://doi.org/10.3390/agriculture11121257

Schlosser, J. F., Farias, M. S. de, Bertollo, G. M., Russini, A., Herzog, D., & Casali, L. (2020). Agricultural tractor engines from the perspective of Agriculture 4.0. *REVISTA CIÊNCIA AGRONÔMICA*, *51*(5). https://doi.org/10.5935/1806-6690.20200094

Shao, D., Ishengoma, F., Anastasija Nikiforova, & Mrisho Swetu. (2024). Comparative analysis of data protection regulations in East African countries. *Digital Policy Regulation and Governance*, *26*(30-35). https://doi.org/10.1108/dprg-06-2024-0120

Ujang Paman, Uchida, S., Inaba, S., & Kojima, T. (2017). CAUSES OF TRACTOR BREAKDOWNS AND REQUISITE SOLUTIONS: A CASE STUDY OF SMALL TRACTOR USE IN RIAU PROVINCE, INDONESIA. *ASEAN Journal on Science & Technology for Development*, *25*(1), 27–36. https://doi.org/10.29037/ajstd.228

Zhang, W., Li, C., Peng, G., Chen, Y., & Zhang, Z. (2018). A deep convolutional neural network with new training methods for bearing fault diagnosis under a noisy environment and different working loads. *Mechanical Systems and Signal Processing*, *100*(31-40), 439–453. https://doi.org/10.1016/j.ymssp.2017.06.022